

**PARALLEL AND DISTRIBUTED**

**COMPUTING  
(CSE4001)**

**CONSOLIDATED LAB REPORT**

Name: O G Ragavi

Reg No: 20BCE1988

Slot: D1

Course code: CSE4001

## CSE4001-Parallel and distributed computing

**Name:** O G RAGAVI

**Reg no:** 20BCE1988

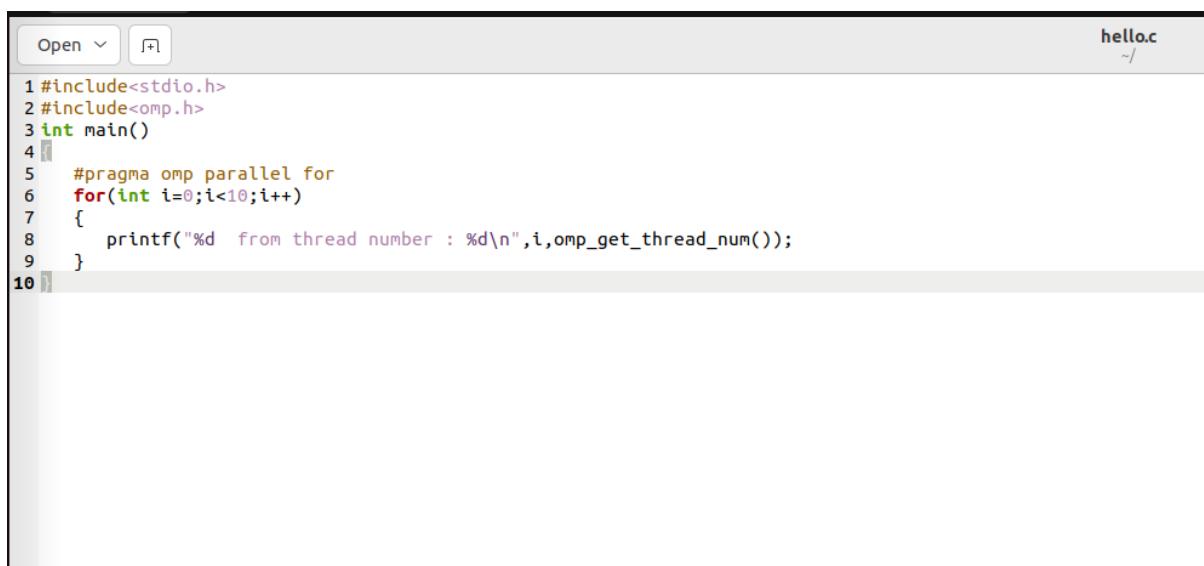
**Lab Ex:** 1-3

**Date:** 21.08.2022

---

### 1. Sample program

**Code:**



The screenshot shows a code editor window with the file 'hello.c' open. The code is a simple C program using OpenMP to print the thread number for each iteration of a loop. The code is as follows:

```
1 #include<stdio.h>
2 #include<omp.h>
3 int main()
4 {
5     #pragma omp parallel for
6     for(int i=0;i<10;i++)
7     {
8         printf("%d from thread number : %d\n",i,omp_get_thread_num());
9     }
10 }
```

**Output:**

The screenshot shows a terminal window with a dark background and light-colored text. At the top right, it says "ragavi@ragavi-VirtualBox: ~". The terminal displays the following command-line session:

```
ragavi@ragavi-VirtualBox:~$ gcc -o hello hello.c
ragavi@ragavi-VirtualBox:~$ ./hello
Hello World
ragavi@ragavi-VirtualBox:~$ gcc -o hello -fopenmp hello.c
ragavi@ragavi-VirtualBox:~$ ./hello
0 from thread number : 0
1 from thread number : 0
2 from thread number : 0
3 from thread number : 0
7 from thread number : 2
8 from thread number : 2
9 from thread number : 2
4 from thread number : 1
5 from thread number : 1
6 from thread number : 1
ragavi@ragavi-VirtualBox:~$ gcc -o hello -fopenmp hello.c
ragavi@ragavi-VirtualBox:~$ ./hello
7 from thread number : 2
8 from thread number : 2
9 from thread number : 2
0 from thread number : 0
1 from thread number : 0
2 from thread number : 0
3 from thread number : 0
4 from thread number : 1
5 from thread number : 1
6 from thread number : 1
ragavi@ragavi-VirtualBox:~$ ss
```

## 2. Program to find the sum of first n numbers.

### Algorithm:

- Start the clock before each and every iteration and stop it before it ends.
- N is initialised and for every iteration no of threads is set.
- Using reduction , we run a for loop and increment the sum by i
- We end the clock and note the time difference(total time taken)

### Code:

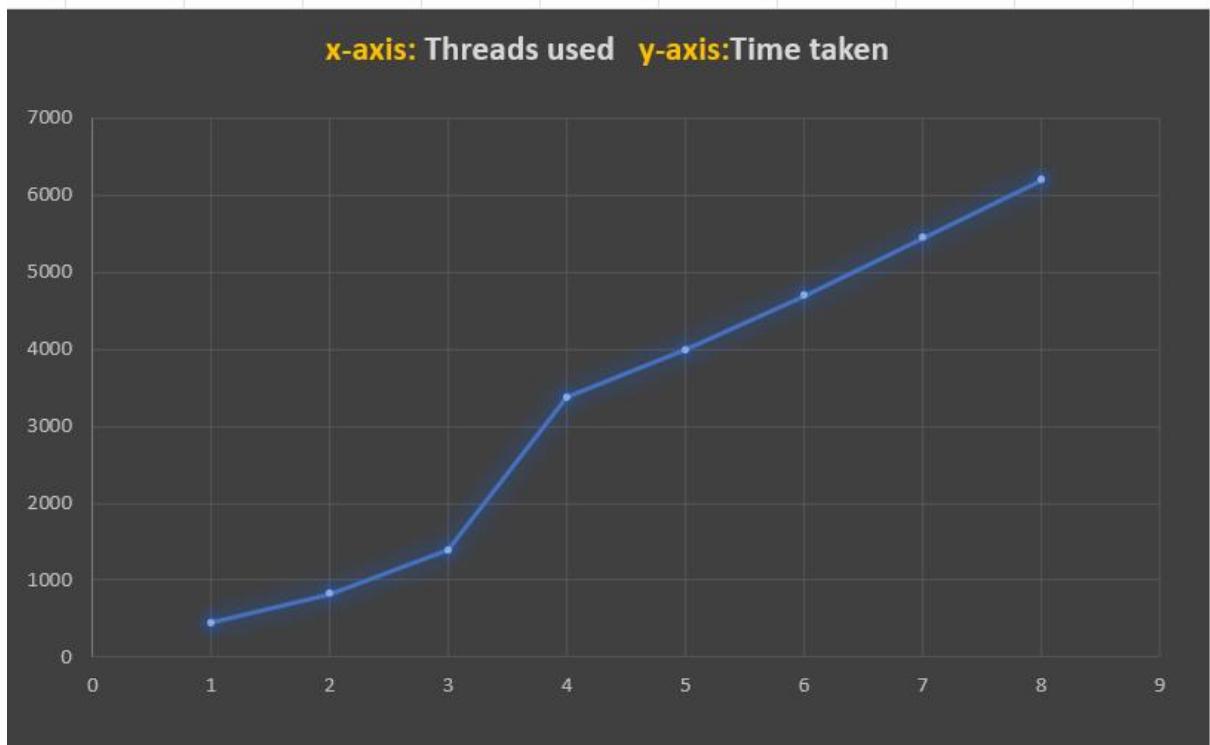
```
Open [+]
prog1.c
~/

1 //program to find the sum of first n numbers where n>10000 with time constraint.
2 #include<stdio.h>
3 #include<omp.h>
4 #include<time.h>
5 int main()
6 {
7     clock_t s,e;
8
9     int i;
10    int n=10000,sum=0;
11    for(int j=1;j<=8;j++)
12    {
13        s=clock();
14        omp_set_num_threads(j); // In which thread the corresponding function is working
15        #pragma omp parallel for reduction(+: sum)
16        for(i=0;i<n;i++)
17        {
18            sum=sum+i;
19        }
20
21        printf("Sum is : %d",sum);
22        e=clock();
23        printf("\nTime taken : %ld for thread : %d\n",e-s,j);
24        sum=0;
25    }
26 }
```

## Output:

```
[+]
ragavi@ragavi-VirtualBox:~$ gcc -o prog1 -fopenmp prog1.c
ragavi@ragavi-VirtualBox:~$ ./prog1
Sum is : 704982704
Time taken : 373 Sum is : 704982704
Time taken : 716 Sum is : 704982704
Time taken : 978 Sum is : 704982704
Time taken : 2272 Sum is : 704982704
Time taken : 2795 Sum is : 704982704
Time taken : 3402 Sum is : 704982704
Time taken : 3955 Sum is : 704982704
Time taken : 4961 ragavi@ragavi-VirtualBox:~$ gcc -o prog1 -fopenmp prog1.c
ragavi@ragavi-VirtualBox:~$ ./prog1
Sum is : 704982704
Time taken : 442 for thread : 1
Sum is : 704982704
Time taken : 821 for thread : 2
Sum is : 704982704
Time taken : 1392 for thread : 3
Sum is : 704982704
Time taken : 3377 for thread : 4
Sum is : 704982704
Time taken : 3987 for thread : 5
Sum is : 704982704
Time taken : 4696 for thread : 6
Sum is : 704982704
Time taken : 5440 for thread : 7
Sum is : 704982704
Time taken : 6196 for thread : 8
ragavi@ragavi-VirtualBox:~$
```

## Graph Plot:



## 2.Program to find the sum of n numbers: n>10000

### Algorithm:

- Start the clock before each and every iteration and stop it before it ends.
- N is initialised and for every iteration no of threads is set.
- Using reduction , we run a for loop and increment the sum by any random number generated.(0-100)
- We end the clock and note the time difference(total time taken)

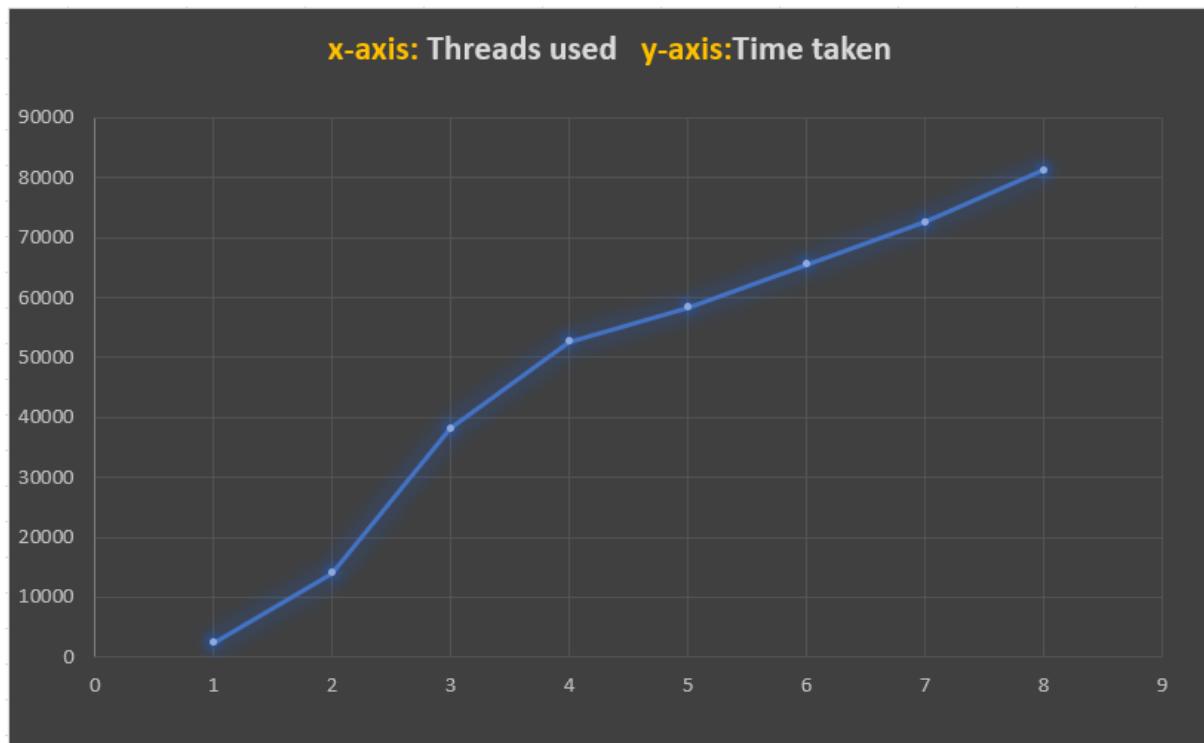
### Code:

```
Open ▾  prog2.c
~/
1 //program to find the sum of n numbers where n>10000 with time constraint.
2 #include<stdio.h>
3 #include<omp.h>
4 #include<time.h>
5 #include<stdlib.h>
6 int main()
7 {
8     clock_t s,e;
9
10    int i;
11    int n=100001,sum=0;
12    for(int j=1;j<=8;j++)
13    {
14        s=clock();
15        omp_set_num_threads(j); // In which thread the corresponding function is working
16        #pragma omp parallel for reduction(+: sum)
17        for(i=0;i<n;i++)
18        {
19            sum=sum+(rand()%100);
20        }
21    }
22
23    printf("Sum is : %d",sum);
24    e=clock();
25    printf("\nTime taken : %ld for thread : %d\n",e-s,j);
26    sum=0;
27 }
28
```

## Output:

```
ragavi@ragavi-VirtualBox:~/Documents$ ./prog2
Sum is : 4952488
Time taken : 2436 for thread : 1
Sum is : 4949527
Time taken : 14468 for thread : 2
Sum is : 4942504
Time taken : 43845 for thread : 3
Sum is : 4948385
Time taken : 62964 for thread : 4
Sum is : 4953382
Time taken : 66958 for thread : 5
Sum is : 4953961
Time taken : 90371 for thread : 6
Sum is : 4960295
Time taken : 94243 for thread : 7
Sum is : 4940623
Time taken : 97459 for thread : 8
```

## Graph Plot:



### 3. Program to find the addition of 2 integer arrays :

#### Algorithm:

- Initialise 2 arrays a and b.
- Use a for loop to get the sum at each index.
- Omp\_get\_num\_threads – building function that takes void input and returns no threads.
- Stop the clock.

#### Code:

```
Open ▾  pro3.c ~/  
1 //program to find the addition of 2 integer arrays with time constraint.  
2 #include<stdio.h>  
3 #include<omp.h>  
4 #include<time.h>  
5 #include<stdlib.h>  
6 int main()  
7 {  
8     clock_t s,e;  
9  
10    int i;  
11    int tid;  
12    int a[5]={0,1,2,3,4};  
13    int b[5]={0,1,2,3,4};  
14    int res[5];  
15    for(int j=1;j<=8;j++)  
16    {  
17        s=clock();  
18        #pragma omp parallel for schedule(static) num_threads(j-1)  
19        for(i=0;i<5;i++)  
20        {  
21            res[i]=a[i]+b[i];  
22            printf("Thread number:%d,res[%d]= %d\n", omp_get_thread_num(),i,res[i]);  
23        }  
24  
25  
26  
27    printf("\n");  
28    e=clock();  
29    printf("\nTime taken : %ld for no of threads : %d\n",e-s,j);  
30  
31 }  
32 }
```

## Output:

```
ragavi@ragavi-VirtualBox:~$ ./pro3
Thread number:0,res[0]= 0
Thread number:0,res[1]= 2
Thread number:1,res[2]= 4
Thread number:1,res[3]= 6
Thread number:2,res[4]= 8

Time taken : 1830 for no of threads : 1
Thread number:0,res[0]= 0
Thread number:0,res[1]= 2
Thread number:0,res[2]= 4
Thread number:0,res[3]= 6
Thread number:0,res[4]= 8

Time taken : 7 for no of threads : 2
Thread number:0,res[0]= 0
Thread number:0,res[1]= 2
Thread number:0,res[2]= 4
Thread number:1,res[3]= 6
Thread number:1,res[4]= 8

Time taken : 761 for no of threads : 3
Thread number:0,res[0]= 0
Thread number:0,res[1]= 2
Thread number:2,res[4]= 8
Thread number:1,res[2]= 4
Thread number:1,res[3]= 6

Time taken : 177 for no of threads : 4
Thread number:2,res[3]= 6
Thread number:1,res[2]= 4
Thread number:0,res[0]= 0
Thread number:0,res[1]= 2
Thread number:3,res[4]= 8

Time taken : 191 for no of threads : 5
Thread number:1,res[1]= 2
Thread number:4,res[4]= 8
Thread number:3,res[3]= 6
Thread number:0,res[0]= 0
Thread number:2,res[2]= 4
```

```
[root]                                     ragavi@ragavi-VirtualBox: ~
Thread number:0,res[1]= 2
Thread number:0,res[2]= 4
Thread number:1,res[3]= 6
Thread number:1,res[4]= 8

Time taken : 761 for no of threads : 3
Thread number:0,res[0]= 0
Thread number:0,res[1]= 2
Thread number:2,res[4]= 8
Thread number:1,res[2]= 4
Thread number:1,res[3]= 6

Time taken : 177 for no of threads : 4
Thread number:2,res[3]= 6
Thread number:1,res[2]= 4
Thread number:0,res[0]= 0
Thread number:0,res[1]= 2
Thread number:3,res[4]= 8

Time taken : 191 for no of threads : 5
Thread number:1,res[1]= 2
Thread number:4,res[4]= 8
Thread number:3,res[3]= 6
Thread number:0,res[0]= 0
Thread number:2,res[2]= 4

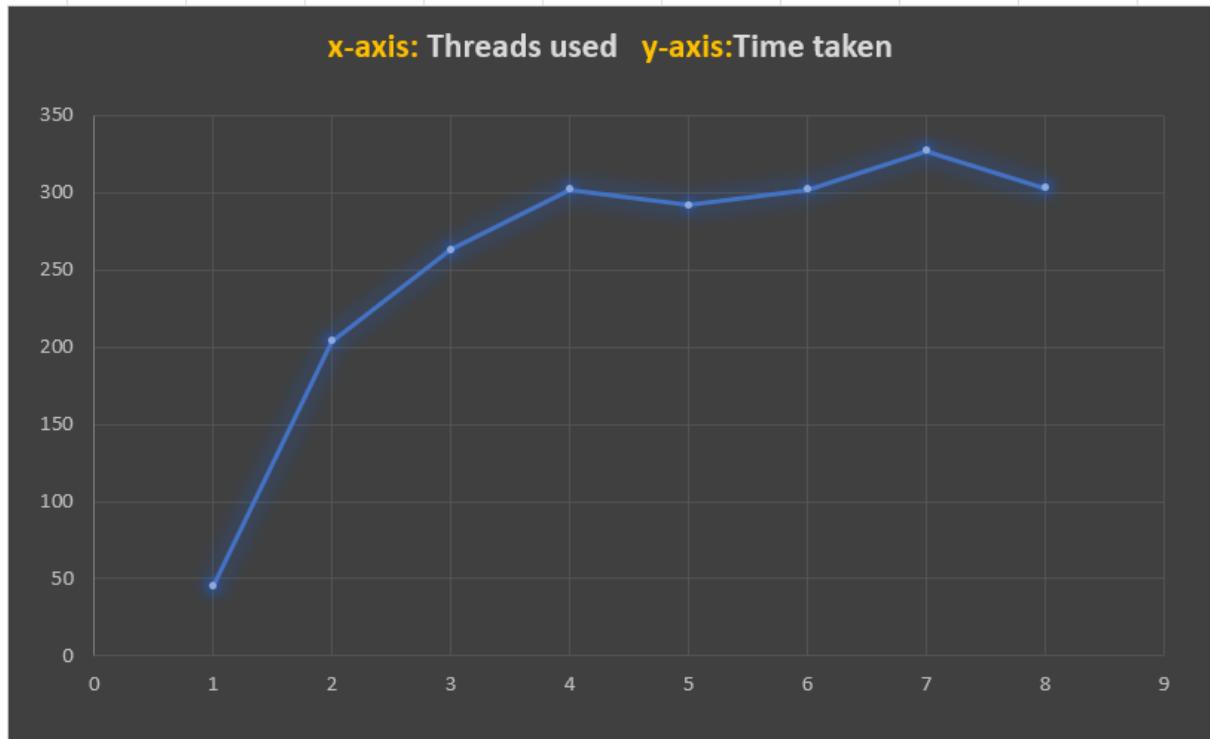
Time taken : 161 for no of threads : 6
Thread number:1,res[1]= 2
Thread number:0,res[0]= 0
Thread number:4,res[4]= 8
Thread number:3,res[3]= 6
Thread number:2,res[2]= 4

Time taken : 160 for no of threads : 7
Thread number:3,res[3]= 6
Thread number:2,res[2]= 4
Thread number:0,res[0]= 0
Thread number:1,res[1]= 2
Thread number:4,res[4]= 8

Time taken : 202 for no of threads : 8
ragavi@ragavi-VirtualBox: ~
```

### Graph Plot:

x axis	y axis
Threads used	Time(s)
1	45
2	204
3	263
4	302
5	292
6	302
7	327
8	303



**4. Program to find the maximum element from an array of 1000 numbers.**

**Code:**

```
Open ▾ prog4.c
~/
1 //program to find the maximum element from an array of 1000 numbers with time constraint.
2 #include<stdio.h>
3 #include<omp.h>
4 #include<time.h>
5 #include<stdlib.h>
6 int main()
7 {
8     clock_t s,e;
9
10    int i,max=0;
11    int tid,a[1000];
12    for(int k=0;k<1000;k++)
13    {
14        a[k]=rand()%50;
15    }
16    for(int j=1;j<=8;j++)
17    {
18        s=clock();
19        #pragma omp parallel for schedule(static) num_threads(j-1)
20        for(i=0;i<1000;i++)
21        {
22            if(max<a[i])
23                max=a[i];
24        }
25        printf("Thread number:%d, max= %d\n", omp_get_thread_num(),max);
26        max=0;
27    }
28
29
30    printf("\n");
31    e=clock();
32    printf("\nTime taken : %ld for no of threads : %d\n",e-s,j);
33
34 }
35 }
36 }
```

## Output:

```
Time taken : 202 for no of threads : 0
ragavi@ragavi-VirtualBox:~$ gcc -o prog4 -fopenmp prog4.c
ragavi@ragavi-VirtualBox:~$ ./prog4
Thread number:0, max= 49

Time taken : 198 for no of threads : 1
Thread number:0, max= 49

Time taken : 6 for no of threads : 2
Thread number:0, max= 49

Time taken : 6 for no of threads : 3
Thread number:0, max= 49

Time taken : 259 for no of threads : 4
Thread number:0, max= 49

Time taken : 389 for no of threads : 5
Thread number:0, max= 49

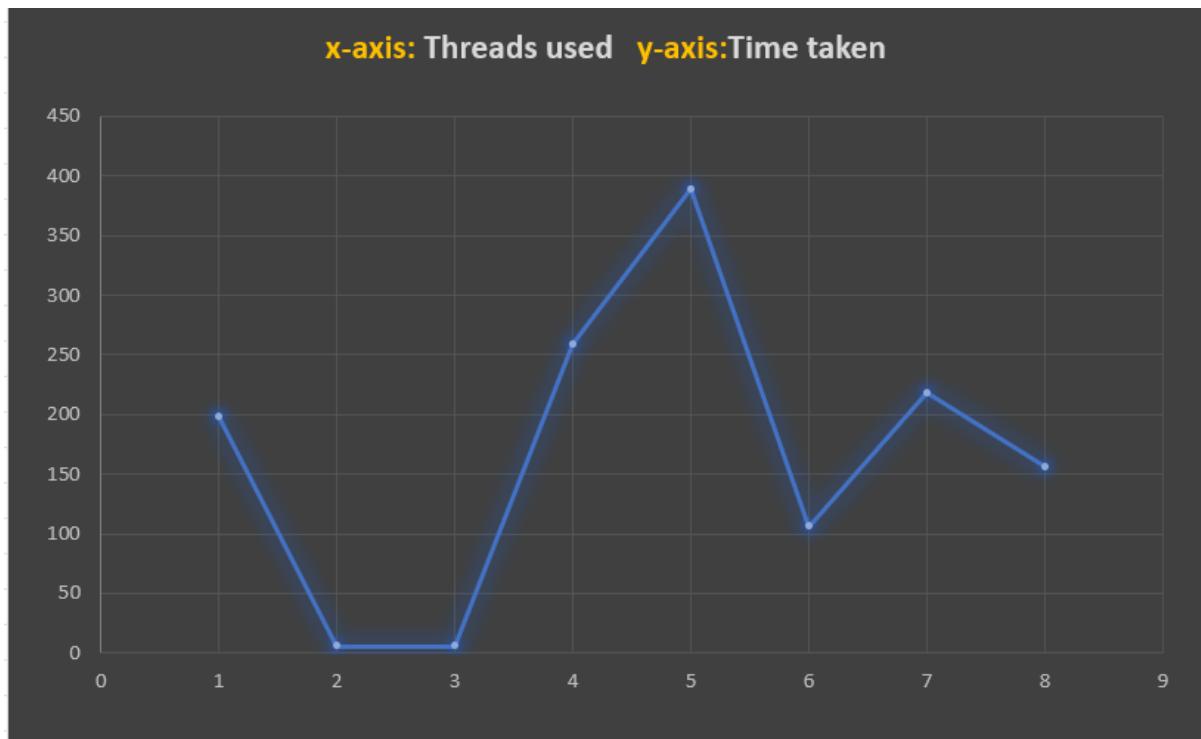
Time taken : 106 for no of threads : 6
Thread number:0, max= 49

Time taken : 218 for no of threads : 7
Thread number:0, max= 49

Time taken : 156 for no of threads : 8
ragavi@ragavi-VirtualBox:~$
```

### Graph Plot:

x axis	y axis
Threads used	Time(s)
1	198
2	6
3	6
4	259
5	389
6	106
7	218
8	156



## 5.Program to find the minimum element from the array of n numbers.

**Code:**

```

Open ▾ ⌂ prog5.c ~/
1 //program to find the minimum element from an array of 10000 numbers with time constraint.
2 #include<stdio.h>
3 #include<omp.h>
4 #include<time.h>
5 #include<stdlib.h>
6 int main()
7 {
8     clock_t s,e;
9
10    int i,min1=1000000;
11    int tid,a[10000];
12    for(int k=0;k<10000;k++)
13    {
14        a[k]=rand()%100;
15    }
16    for(int j=1;j<=8;j++)
17    {
18        s=clock();
19        #pragma omp parallel reduction(min:min1)
20        #pragma omp for
21        for(i=0;i<1000;i++)
22        {
23            if(min1>a[i])
24                min1=a[i];
25        }
26    }
27    printf("Thread number:%d, min= %d\n", omp_get_thread_num(),min1);
28    min1=100000;
29
30
31
32    printf("\n");
33    e=clock();
34    printf("\nTime taken : %ld for no of threads : %d\n",e-s,j);
35
36 }
37

```

**Output:**

```
ragavi@ragavi-VirtualBox:~$ gcc -o prog5 -fopenmp prog5.c
ragavi@ragavi-VirtualBox:~$ ./prog5
Thread number:0, min= 0

Time taken : 206 for no of threads : 1
Thread number:0, min= 0

Time taken : 7 for no of threads : 2
Thread number:0, min= 0

Time taken : 5 for no of threads : 3
Thread number:0, min= 0

Time taken : 5 for no of threads : 4
Thread number:0, min= 0

Time taken : 6 for no of threads : 5
Thread number:0, min= 0

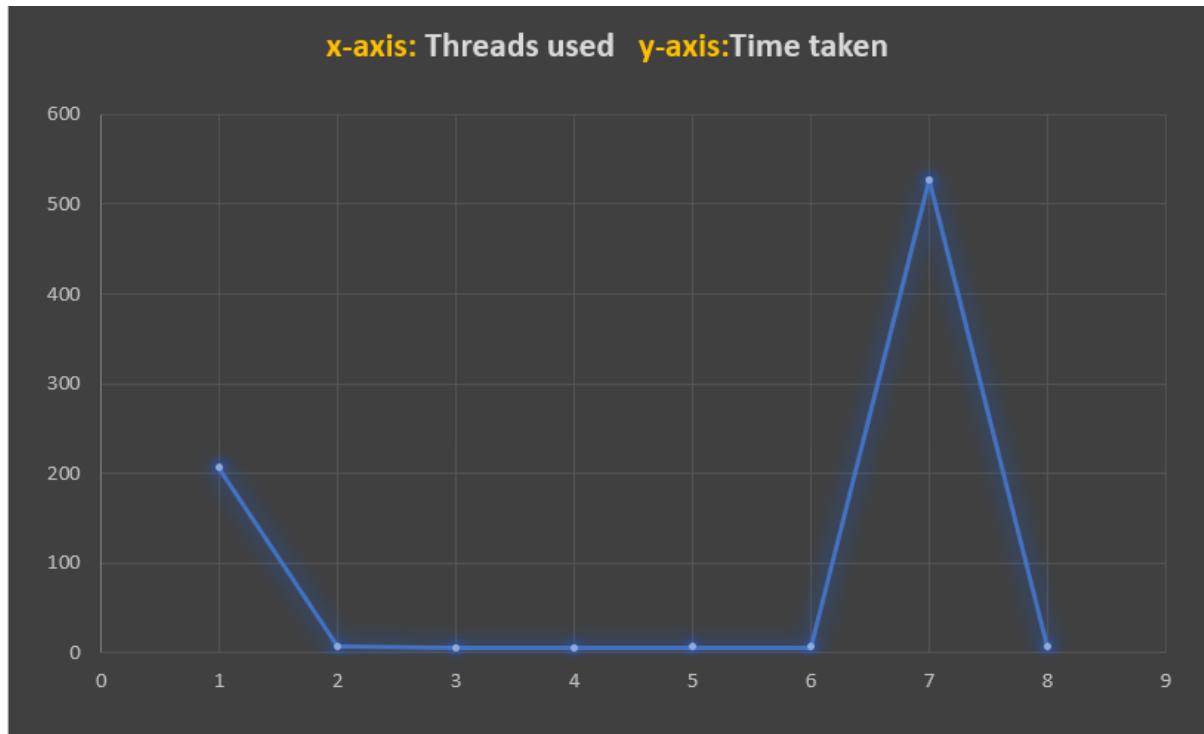
Time taken : 6 for no of threads : 6
Thread number:0, min= 0

Time taken : 527 for no of threads : 7
Thread number:0, min= 0

Time taken : 6 for no of threads : 8
```

### Graph Plot:

x axis	y axis
Threads used	Time(s)
1	206
2	7
3	5
4	5
5	6
6	6
7	527
8	6



## 6.Program to find resultant of matrix multiplication.

**Code:**

```

Open ▾ Run prog6.c ~
1 //program to perform matrix multiplication with time constraint.
2 #include<stdio.h>
3 #include<omp.h>
4 #include<time.h>
5 #include<stdlib.h>
6 int main()
7 {
8     clock_t s,e;
9
10    int i;
11    int tid;
12    int a[3][3]={{1,2,3},{1,2,3},{1,2,3}};
13    int b[3][3]={{1,2,3},{1,2,3},{1,2,3}};
14    int res[3][3];
15    for(int j=1;j<=8;j++)
16    {
17        s=clock();
18        omp_set_num_threads(j);
19        #pragma omp parallel for collapse(2) schedule(static)
20        for(int i1=0;i1<3;i1++)
21        {
22            for(int j1=0;j1<3;j1++)
23            {
24                res[i1][j1]=0;
25                for(int k=0;k<3;k++)
26                {
27                    res[i1][j1]+= a[i1][k]*b[k][j1];
28                }
29            }
30        }
31        for(i=0;i<3;i++)
32        {
33            for(int k=0;k<3;k++)
34            {
35                printf("res[%d][%d]= %d\n",i,k,res[i][k]);
36            }
37        }
38    }
39
40
41
42
43
44    printf("\n");
45    e=clock();
46    printf("Time taken : %ld for no of threads : %d\n",e-s,j);
47

```

## Output:

---

```
[+]
ragavi@ragavi-VirtualBox:~$ gcc -o prog6 -fopenmp prog6.c
ragavi@ragavi-VirtualBox:~$ ./prog6
res[0][0]= 6
res[0][1]= 12
res[0][2]= 18
res[1][0]= 6
res[1][1]= 12
res[1][2]= 18
res[2][0]= 6
res[2][1]= 12
res[2][2]= 18

Time taken : 43 for no of threads : 1
res[0][0]= 6
res[0][1]= 12
res[0][2]= 18
res[1][0]= 6
res[1][1]= 12
res[1][2]= 18
res[2][0]= 6
res[2][1]= 12
res[2][2]= 18

Time taken : 185 for no of threads : 2
res[0][0]= 6
res[0][1]= 12
res[0][2]= 18
res[1][0]= 6
res[1][1]= 12
res[1][2]= 18
res[2][0]= 6
res[2][1]= 12
res[2][2]= 18

Time taken : 3623 for no of threads : 3
res[0][0]= 6
res[0][1]= 12
res[0][2]= 18
res[1][0]= 6
res[1][1]= 12
res[1][2]= 18
res[2][0]= 6
res[2][1]= 12
res[2][2]= 18

Time taken : 519 for no of threads : 4
res[0][0]= 6
res[0][1]= 12
```

```
Time taken : 3623 for no of threads : 3
res[0][0]= 6
res[0][1]= 12
res[0][2]= 18
res[1][0]= 6
res[1][1]= 12
res[1][2]= 18
res[2][0]= 6
res[2][1]= 12
res[2][2]= 18

Time taken : 519 for no of threads : 4
res[0][0]= 6
res[0][1]= 12
res[0][2]= 18
res[1][0]= 6
res[1][1]= 12
res[1][2]= 18
res[2][0]= 6
res[2][1]= 12
res[2][2]= 18

Time taken : 348 for no of threads : 5
res[0][0]= 6
res[0][1]= 12
res[0][2]= 18
res[1][0]= 6
res[1][1]= 12
res[1][2]= 18
res[2][0]= 6
res[2][1]= 12
res[2][2]= 18

Time taken : 308 for no of threads : 6
res[0][0]= 6
res[0][1]= 12
res[0][2]= 18
res[1][0]= 6
res[1][1]= 12
res[1][2]= 18
res[2][0]= 6
res[2][1]= 12
res[2][2]= 18

Time taken : 273 for no of threads : 7
res[0][0]= 6
res[0][1]= 12
res[0][2]= 18
```

```
[+]

res[2][2]= 18

Time taken : 519 for no of threads : 4
res[0][0]= 6
res[0][1]= 12
res[0][2]= 18
res[1][0]= 6
res[1][1]= 12
res[1][2]= 18
res[2][0]= 6
res[2][1]= 12
res[2][2]= 18

Time taken : 348 for no of threads : 5
res[0][0]= 6
res[0][1]= 12
res[0][2]= 18
res[1][0]= 6
res[1][1]= 12
res[1][2]= 18
res[2][0]= 6
res[2][1]= 12
res[2][2]= 18

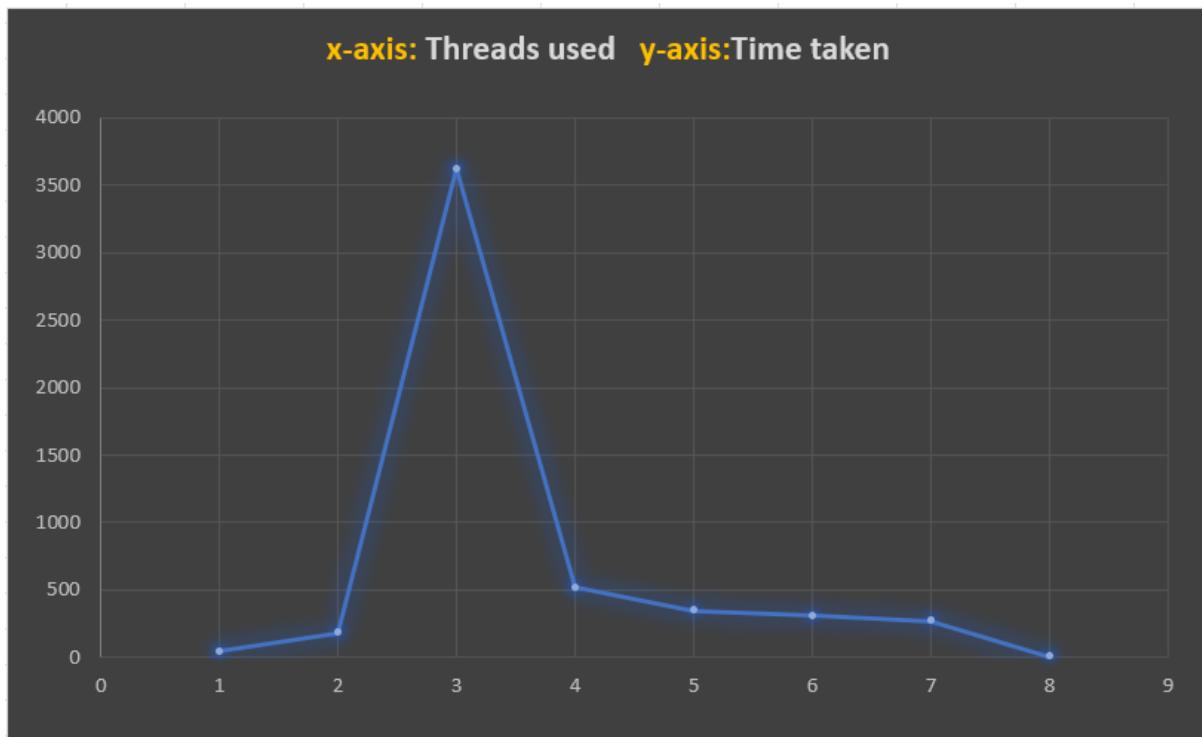
Time taken : 308 for no of threads : 6
res[0][0]= 6
res[0][1]= 12
res[0][2]= 18
res[1][0]= 6
res[1][1]= 12
res[1][2]= 18
res[2][0]= 6
res[2][1]= 12
res[2][2]= 18

Time taken : 273 for no of threads : 7
res[0][0]= 6
res[0][1]= 12
res[0][2]= 18
res[1][0]= 6
res[1][1]= 12
res[1][2]= 18
res[2][0]= 6
res[2][1]= 12
res[2][2]= 18

Time taken : 328 for no of threads : 8
ragavi@ragavi-VirtualBox:~$
```

### Graph Plot:

x axis	y axis
Threads used	Time(s)
1	43
2	185
3	3623
4	519
5	348
6	308
7	273
8	328



7.a)No-wait:

```
#pragma omp parallel
{
    #pragma omp for nowait
    for (i=1; i<n; i++)
        b[i] = (a[i] + a[i-1]) / 2.0;
    #pragma omp for nowait
    for (i=0; i<m; i++)
        y[i] = sqrt(z[i]);
}
```

b)Barrier:

```
#pragma omp parallel default(shared)
{
    #pragma omp for
    for (i=0; i<n; i++)
    {
        #pragma omp parallel shared(i, n)
        {
```

```
        #pragma omp for
        for (j=0; j<n; j++)
            work(i, j);
    }
}
```

# CSE4001-PARALLEL AND DISTRIBUTED COMPUTING

## LAB-4

**Name:** O G Ragavi

**Reg No:** 20BCE1988

**Lab:** 4

---

**1.Aim :**Sorting an array of 1000 numbers

**Algorithm:**

- Include the preprocessor directives
- Use the clock function of time.h library to start the timer before entering into the parallel section
- Use a loop that runs from 1 to size of the array which is 1000 .
- Use the phase #pragma omp parallel for, so that the iterations of the inner loop gets distributed among the given number of threads to reach maximum efficiency .

**Code:**

```
#include<stdio.h>
```

```
#include<omp.h>
```

```
#include<time.
```

```
h>
```

```
#include<stdlib
```

```
.h> int main () {
```

```
int
```

```
A[1000];
```

```
int N =
```

```
1000; int
i=0;
for(int i=0;i<N;i++)
{
A[i]=rand()%50;
}
clock_t
s,e;

for(int k=0;k<5;k++)
{
s=clock();
omp_set_num_threads(k+1);
#pragma omp parallel for default(none), shared(A, N)
for(int i=0;i<N;i++)
{

    for(int j = 0; j < N-1; j++)
    {
if(A[j] > A[j
+ 1])
{ int temp =
A[j]; A[j] =
A[j + 1];
A[j + 1] = temp;
}

```

```
}

}

/*
for(i=0;i<N;i
++) {
//printf("%d ",A[i]);

}*/ e = clock(); printf("\nTime taken for %d threads
is: %ld\n",k+1,(e-s));

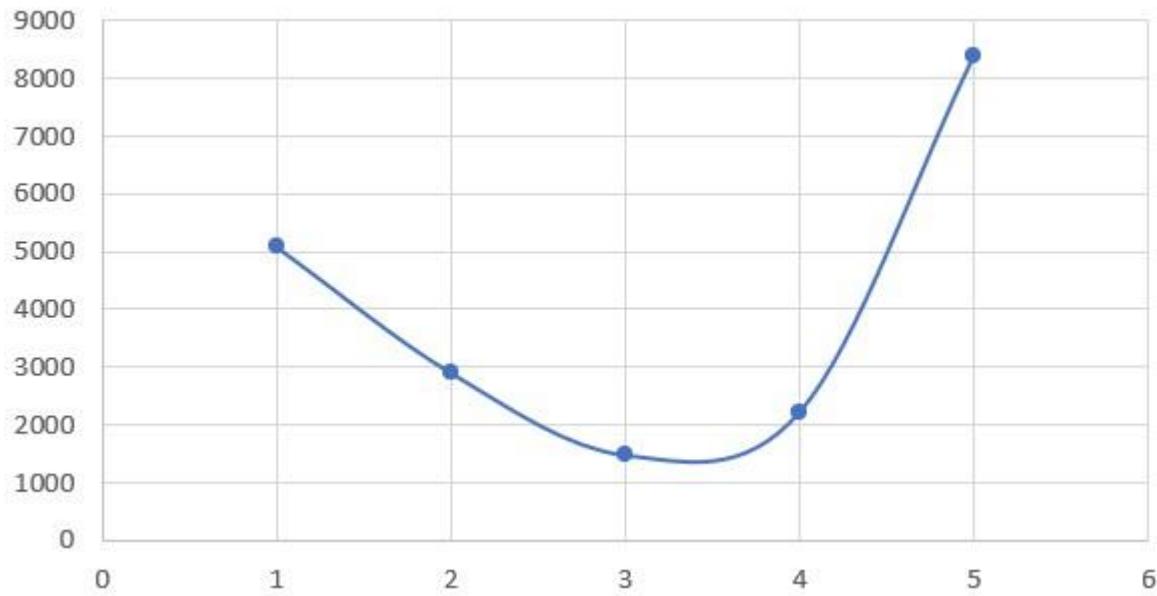
}

}
```

### Output:

```
Time taken for 1 threads is: 5075
Time taken for 2 threads is: 2895
Time taken for 3 threads is: 1475
Time taken for 4 threads is: 2223
Time taken for 5 threads is: 8391
```

### Graph:



## 2.Aim : Searching an element in an array of elements:

### Algorithm:

- Include the required pre-processor directives
- Declare an array of the required size and initialize it with random numbers
- Use the time function to start the clock before the parallel section
- User #pragma omp parallel for, so that the iterations gets distributed among the given number of threads
- Inside the loop use a binary search to search the key from the corresponding array, print the location of the element found.

### Code:

```
#include<stdio.h>
#include<omp.h>
#include<stdlib.h>
#include<time.h>
```

```
int bin(int a[], int l, int r, int key, int n)

{ int index = -1; int
  size = (r - l + 1) /
  n; if (size == 0 ||

n == 1)

{

#pragma omp parallel
for for (int i = l; i <= r;
i++)

{ if (a[i] ==

key)

  index = i;

}

return index;

} int
left = l;
int
right =
r;
omp_set_num_threads(n);
omp_set_nested(1);
#pragma omp parallel
{
  int id = omp_get_thread_num();
  int lt = l + id *
  size; int rt = lt
  + size - 1; if (id
  == n - 1)
```

```

rt = r;

if (a[l] <= key && a[rt] >= key)

{ left =
    lt;
    right
    = rt;

} } if (left == l
&& right == r)

return -1;

return bin(a, left, right, key, n);

}

int main()

{ int n; int nthread; clock_t s, e;
printf("Enter the size of the
array: "); scanf("%d", &n); int
a[n];

for (int i = 0; i < n; i++)

    a[i] = rand()%50; printf("Enter the
    element to be searched: ");

int key;
scanf("%d",
&key); for(int
i=0;i<5;i++){

    s = clock(); printf("Position of the elemnt:%d\n",
    bin(a, 0, n- 1, key, i+1)); e = clock(); printf("Time
    taken for %d threads is : %ld\n",i+1,(e-s));
}

```

```
}
```

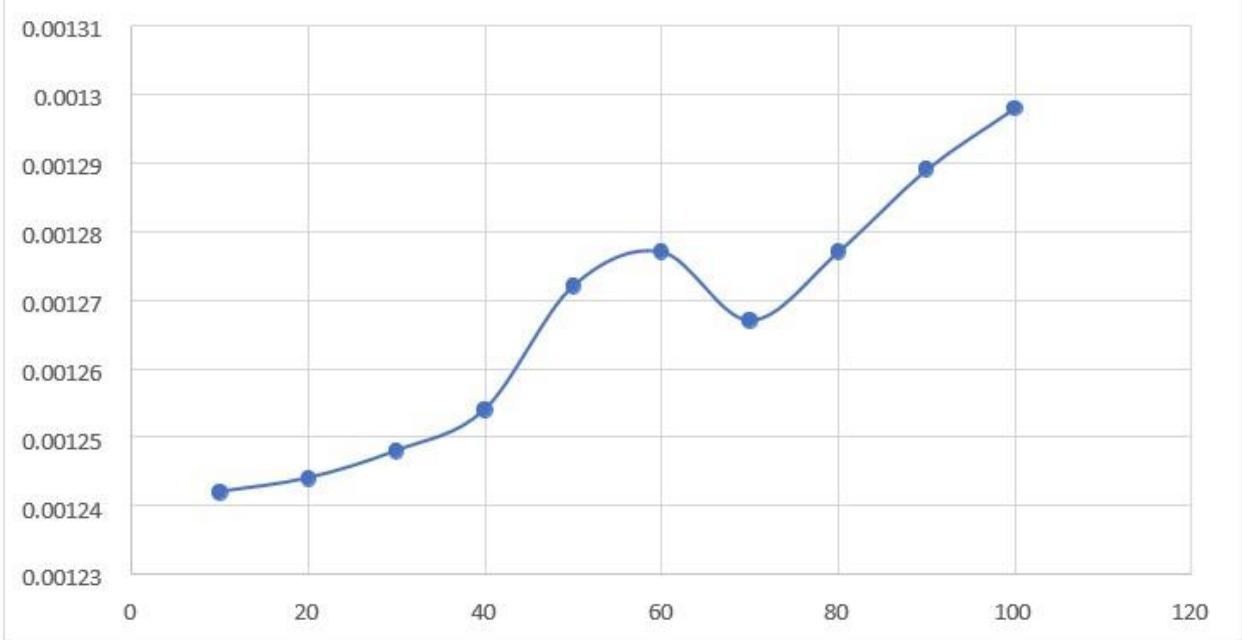
```
}
```

### Output:

```
Enter the size of the array: 100
Enter the element to be searched: 10
Position of the elemnt:91
Time taken for 1 threads is : 347
Position of the elemnt:-1
Time taken for 2 threads is : 11
Position of the elemnt:-1
Terminal n for 3 threads is : 535
Position of the elemnt:-1
Time taken for 4 threads is : 1224
Position of the elemnt:-1
Time taken for 5 threads is : 575
```

```
Enter the size of the array: 100
Enter the element to be searched: 10
Position of the elemnt:91
Time taken for 1 threads is : 347
Position of the elemnt:-1
Time taken for 2 threads is : 11
Position of the elemnt:-1
Time taken for 3 threads is : 535
Position of the elemnt:-1
Time taken for 4 threads is : 1224
Position of the elemnt:-1
Time taken for 5 threads is : 575
```

### Graph:



**3.Aim:** To write an OpenMP program to illustrate master, single, firstprivate, lastprivate and ordered . **a. Master Code:**

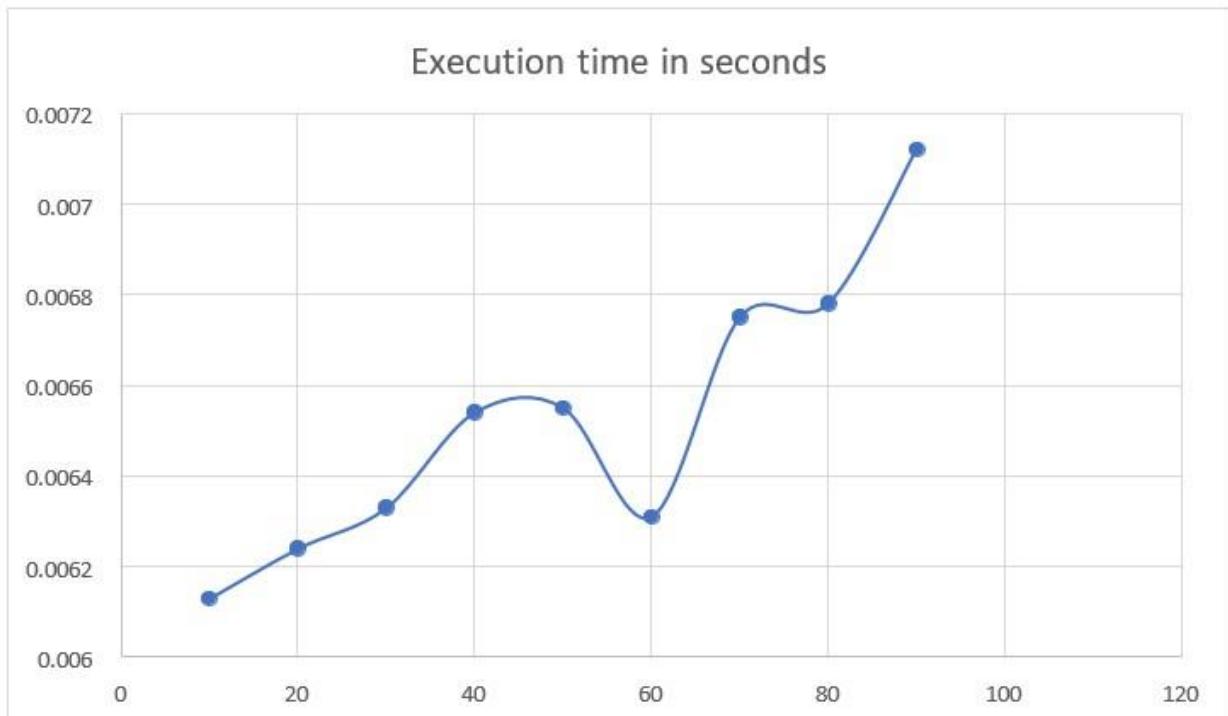
```
#include<stdio.h>
#include<omp.h>
#include<stdlib.h>
#include<time.h>
void main()
{
    time_t
    start,end; start
    = clock();
    #pragma omp parallel
    {
        printf("The block of code is executed by thread
%d\n",omp_get_thread_num());
        #pragma omp master
```

```
{  
    printf("This is the master block and executed by thread  
%d\n",omp_get_thread_num());  
}  
}  
end = clock();  
  
double diff = (end-  
start)*1.0/CLOCKS_PER_SEC;  
printf("The execution time is %lf seconds\n",diff);  
}
```

### Output:

```
The block of code is executed by thread 0  
This is the master block and executed by thread 0  
The block of code is executed by thread 3  
The block of code is executed by thread 4  
The block of code is executed by thread 5  
The block of code is executed by thread 6  
The block of code is executed by thread 7  
The block of code is executed by thread 8  
The block of code is executed by thread 9  
The block of code is executed by thread 2  
The block of code is executed by thread 1  
The execution time is 0.000613 seconds
```

### Graph:



**b.Single:**

**Code:**

```
#include<stdio.h>

#include<omp.h>

#include<stdlib.h>

#include<time.h>
> void main()
{
    time_t
    start,end; start =
    clock();

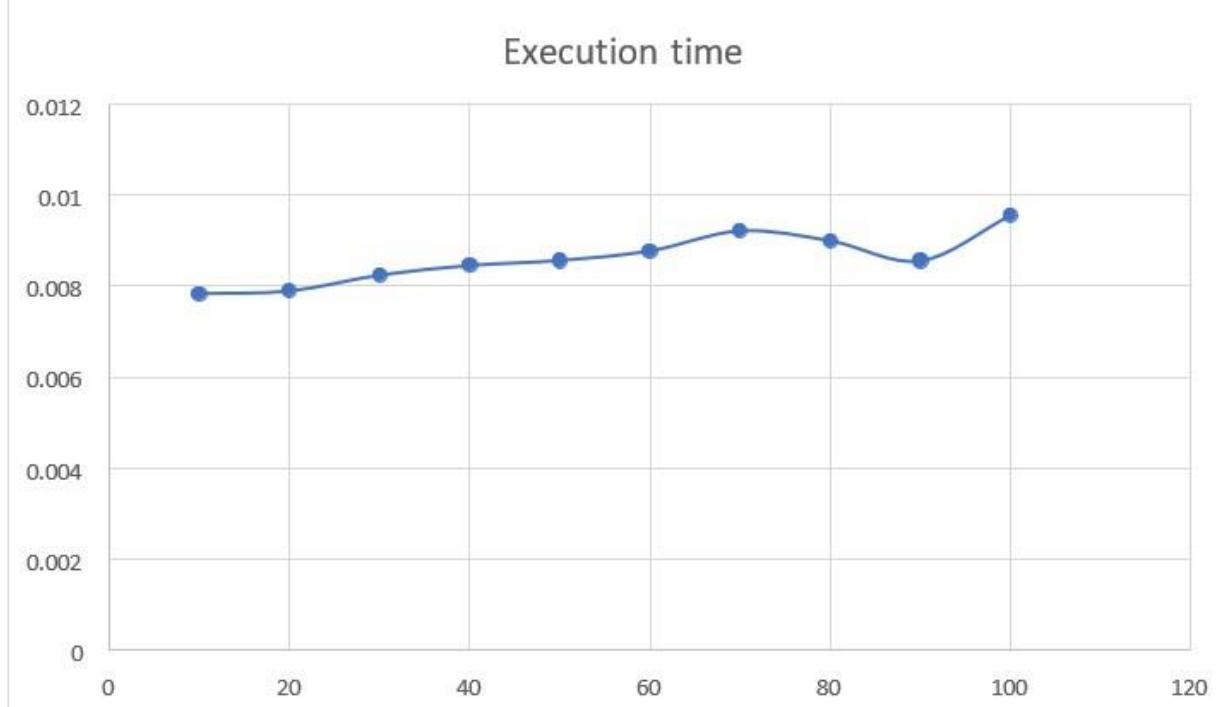
    #pragma omp parallel
    {
        printf("The block of code is executed by thread
%d\n",omp_get_thread_num());      #pragma omp single
        {

```

```
    printf("This is the single block and executed by thread  
%d\n",omp_get_thread_num());  
}  
}  
end = clock();  
double diff = (end-start)*1.0/CLOCKS_PER_SEC;  printf("The  
execution time  
is %lf  
seconds\n",diff);  
}
```

### Output:

```
The block of code is executed by thread 9  
This is the single block and executed by thread 9  
The block of code is executed by thread 1  
The block of code is executed by thread 2  
The block of code is executed by thread 3  
The block of code is executed by thread 4  
The block of code is executed by thread 5  
The block of code is executed by thread 6  
The block of code is executed by thread 7  
The block of code is executed by thread 8  
The block of code is executed by thread 0  
The execution time is 0.000783 seconds
```



### c.Firstprivat

e: Code:

```
#include<stdio.h>

#include<omp.h>

#include<stdlib.h>
#include<time.h>
> void main()
{
    time_t
    start,end;
    int i = 127;
    start =
    clock();
    #pragma omp parallel firstprivate(i)
    {
```

```

        printf("The value of i in the thread %d is
                %d\n",omp_get_thread_num(),i);
    }

end = clock(); double diff = (end-
start)*1.0/CLOCKS_PER_SEC;

printf("The execution time is %lf seconds\n",diff);

}

```

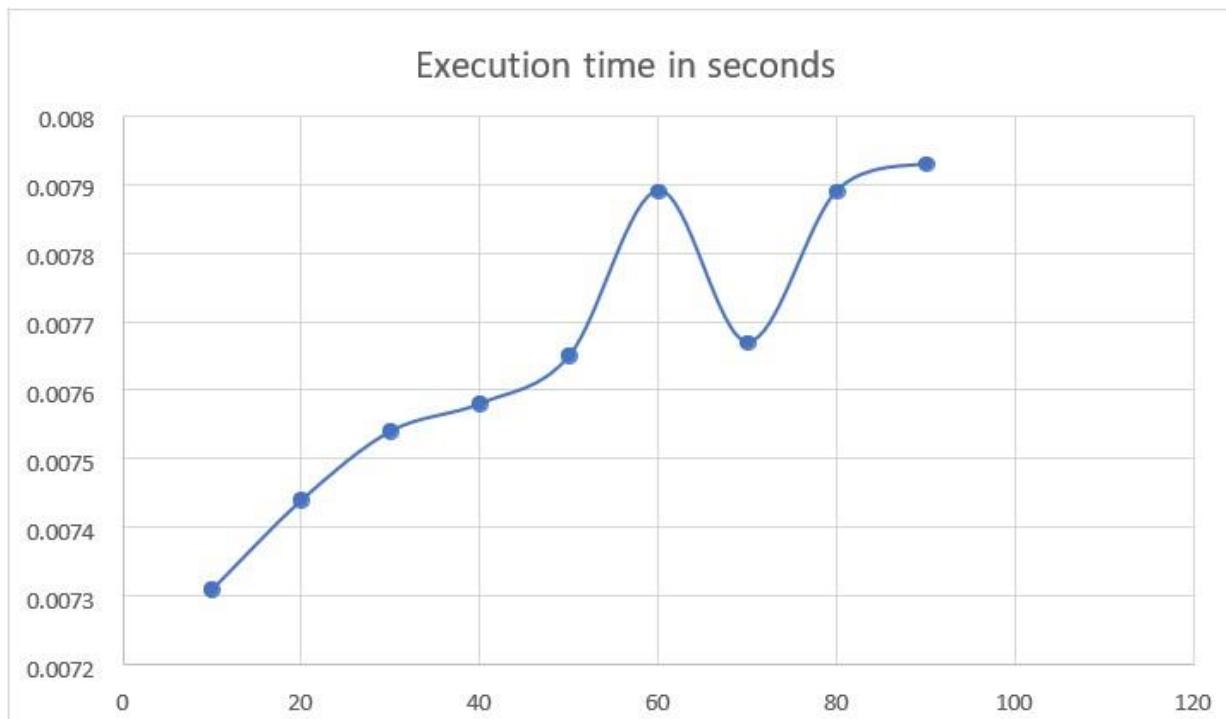
**Output:**

```

The value of i in the thread 9 is 127
The value of i in the thread 0 is 127
The value of i in the thread 1 is 127
The value of i in the thread 2 is 127
The value of i in the thread 3 is 127
The value of i in the thread 4 is 127
The value of i in the thread 5 is 127
The value of i in the thread 6 is 127
The value of i in the thread 7 is 127
The value of i in the thread 8 is 127
The execution time is 0.000731 seconds

```

**Graph:**



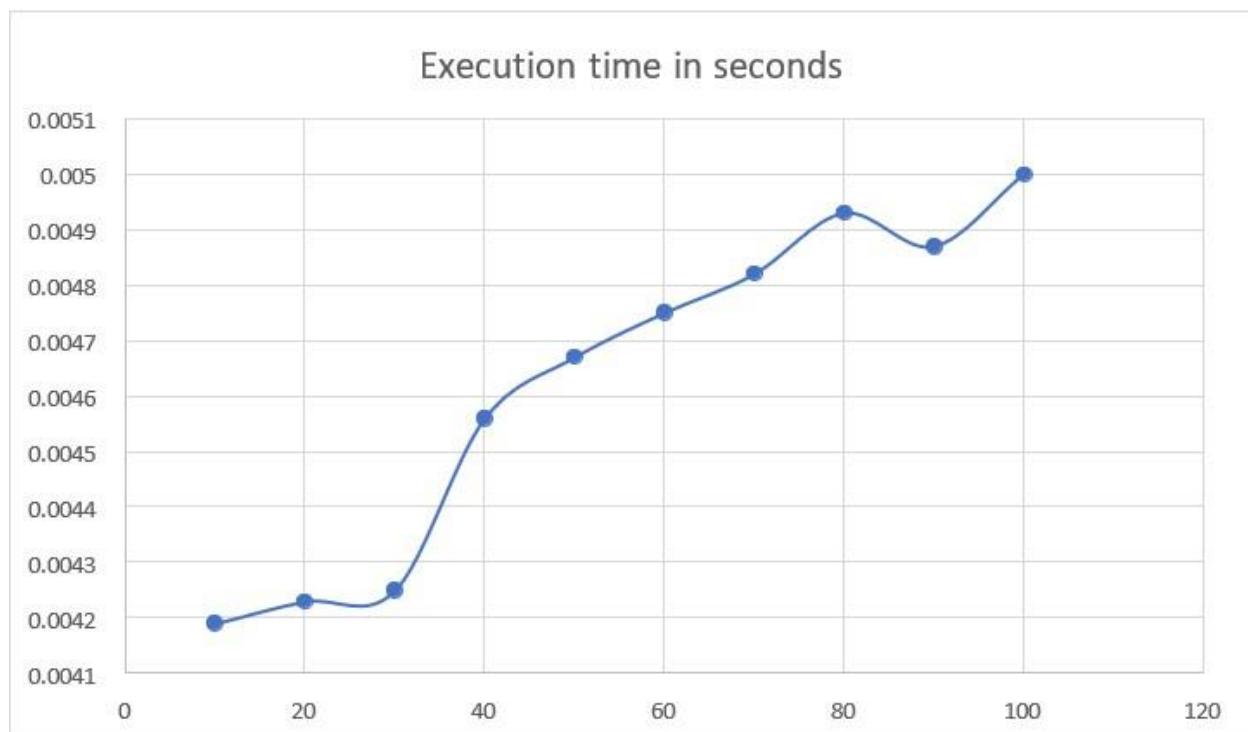
**d.Lastprivat****e: Code:**

```
#include<stdio.h>
#include<omp.h>
#include<stdlib.h>
#include<time.h>
> void main()
{
    time_t
start,end; start
= clock(); int j =
    0;
#pragma omp parallel for
    lastprivate(j) for(int i=0;i<10;i++)
    {
        printf("The value of j in the thread %d is
            %d\n",omp_get_thread_num(),j);
        j++;
    } printf("The value of j outside the parallel
region is %d\n",j);
end = clock();
    double diff = (end-start)*1.0/CLOCKS_PER_SEC;
printf("The execution time is %lf seconds\n",diff);
}
```

**Output:**

```
The value of j in the thread 3 is 0
The value of j in the thread 3 is 1
The value of j in the thread 1 is 0
The value of j in the thread 1 is 1
The value of j in the thread 1 is 2
The value of j in the thread 2 is 0
The value of j in the thread 2 is 1
The value of j in the thread 0 is 0
The value of j in the thread 0 is 1
The value of j in the thread 0 is 2
The value of j outside the parallel region is 2
The execution time is 0.000419 seconds
```

### Graph:



### D.ordered:

#### Code:

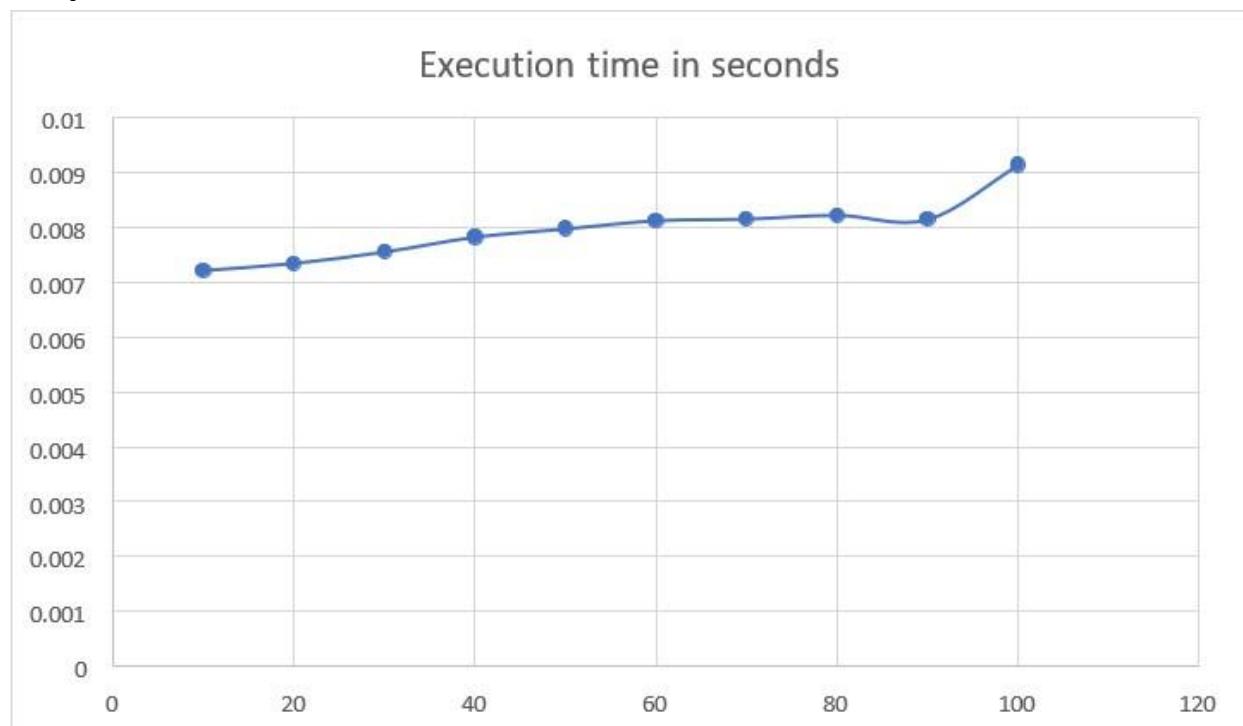
```
#include<stdio.h>
#include<omp.h>
#include<stdlib.h>
#include<time.h>
> void main()
```

```
{\n    time_t\n    start,end;\n    start =\n    clock(); int j\n    = 0;\n\n    #pragma omp parallel\n{\n        int x\n        = 0;\n\n        #pragma omp for\n        ordered for(int\n        i=0;i<10;i++)\n\n        { x =\n            x+i;\n\n            #pragma omp ordered printf("in i=%d x=%d\n            thread=%d\\n",i,x,omp_get_thread_num());\n        }\n    }\n\n    end = clock(); double diff = (end-\n    start)*1.0/CLOCKS_PER_SEC;\n\n    printf("The execution time is %lf seconds\\n",diff);\n}
```

**Output:**

```
in i=0 x=0 thread=0
in i=1 x=1 thread=1
in i=2 x=2 thread=2
in i=3 x=3 thread=3
in i=4 x=4 thread=4
in i=5 x=5 thread=5
in i=6 x=6 thread=6
in i=7 x=7 thread=7
in i=8 x=8 thread=8
in i=9 x=9 thread=9
The execution time is 0.000721 seconds
```

### Graph:



## **CSE4001-PARALLEL AND DISTRIBUTED COMPUTING**

**Name:** O G Ragavi

**Reg No:** 20BCE1988

**Date:** 9-9-2022

**Lab Ex:** 5

---

### **1.Finding the prime numbers in the given range:(for $2^1$ to $2^{17}$ )**

#### **Algorithm:**

- We use a simple sieve algorithm to find the prime numbers in a given range.
- We run a loop for 1-5 threads ; each time we set the number of threads from 1-5 and run the function to find the number of prime numbers within the range.
- Start the clock before setting the thread to the corresponding number and close the clock after printing the count value.

#### **Code:**

```
#include<stdio.h>
#include<omp.h>
#include<stdlib.h>
#include<limits.h>
#include<time.h>
> int main()
{
```

```
int
prime[1000000];
long long i, j, n;
clock_t start,end;
double
cpu_time_used;

printf("\nEnter the value of n");

scanf("%lld",&n);
for(int
k=0;k<5;k++) {
start=clock();
omp_set_num_thr
eads(k+1);

for(i=1;i<=n;i++)
{
prime[i]=1;
}

prime[1]=0;
for(i=2;i*i<=n;i
++) {

#pragma omp parallel for
for(j=i*i;j<=n;j=j+i)

{
if(prime[j]
==1)

prime[j]=0;
```

```
    }

}

int count = 0;
for(i=2;i<=n;i++)
{
if(prime[i] ==
1)
{
    count++;
}
} printf("%d\n",
count);
end=clock();
cpu_time_used=(
(double)(end-
start))/CLOCKS_P
ER_SEC;
printf("Execution
Time is:%lf for
%d
threads\n",cpu_t
ime_used,k+1);
```

```
}
```

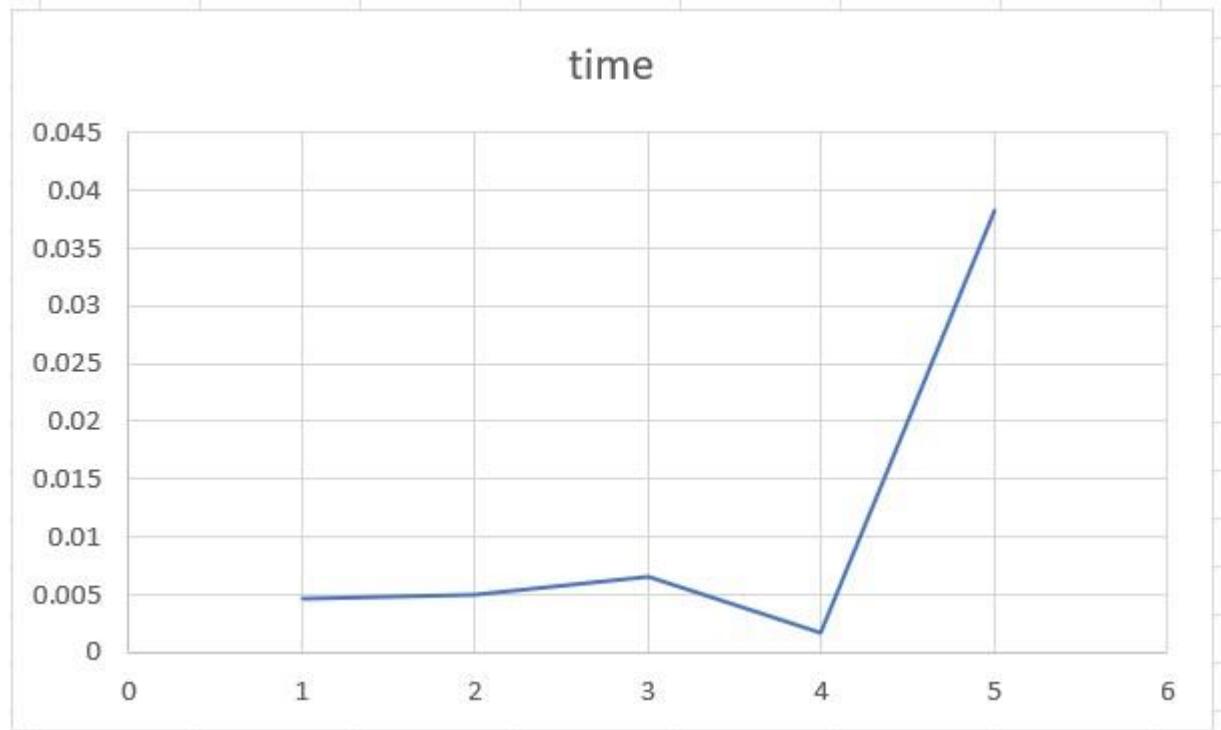
```
}
```

### Output:

```
Enter the value of n 131072
12251
Execution Time is:0.004724 for 1 threads
12251
Execution Time is:0.005071 for 2 threads
12251
Execution Time is:0.006595 for 3 threads
12251
Execution Time is:0.001666 for 4 threads
12251
Execution Time is:0.038366 for 5 threads
```

### Graph:

n=131072	
No of threads	time
1	0.004724
2	0.005071
3	0.006595
4	0.001666
5	0.038366



**CSE4001-PARALLEL AND DISTRIBUTED  
COMPUTING**  
**lab-6**

**Name:** O G Ragavi

**Reg No:** 20BCE1988

**Lab:** 6

---

---

**1.OMP program to multiply 2 matrices:**

**Code:**

```
// C program to multiply two
square matrices.

#include <stdio.h>
#include<math.h>
#include<omp.h>
#include<time.h>
#define N 4

// This function multiplies mat1[][]
and mat2[][][],
// and stores the result in res[][]

void multiply(int mat1[][N], int
mat2[][N], int res[][N])
{ int
    i, j,
    k;
    #pragma omp parallel for
    collapse(2) schedule(static) for (i
= 0; i < N; i++) { for (j = 0; j < N;
j++) { res[i][j] = 0; for (k = 0; k <
```

```

N; k++) res[i][j] += mat1[i][k] *
mat2[k][j]; }
}

int main()
{ clock_t
s,e;
s=clock()
;
int mat1[N][N] = { { 1, 1, 1, 1 },
{ 2, 2, 2, 2 },
{ 3, 3, 3, 3 },
{ 4, 4, 4, 4 } };

int mat2[N][N] = { { 1, 1, 1, 1 },
{ 2, 2, 2, 2 },
{ 3, 3, 3, 3 },
{ 4, 4, 4, 4 } };

int res[N][N]; // To
store result int i, j;
multiply(mat1, mat2, res);

printf("Result
matrix is \n");
for (i = 0; i <
N; i++) { for (j
= 0; j < N; j++)
printf("%d ", res[i][j]);
printf("\n");
}
e=clock();
printf("Time

```

```

    taken %ld ",(e-
s));
return 0;
}

```

```

matrix_mult.c

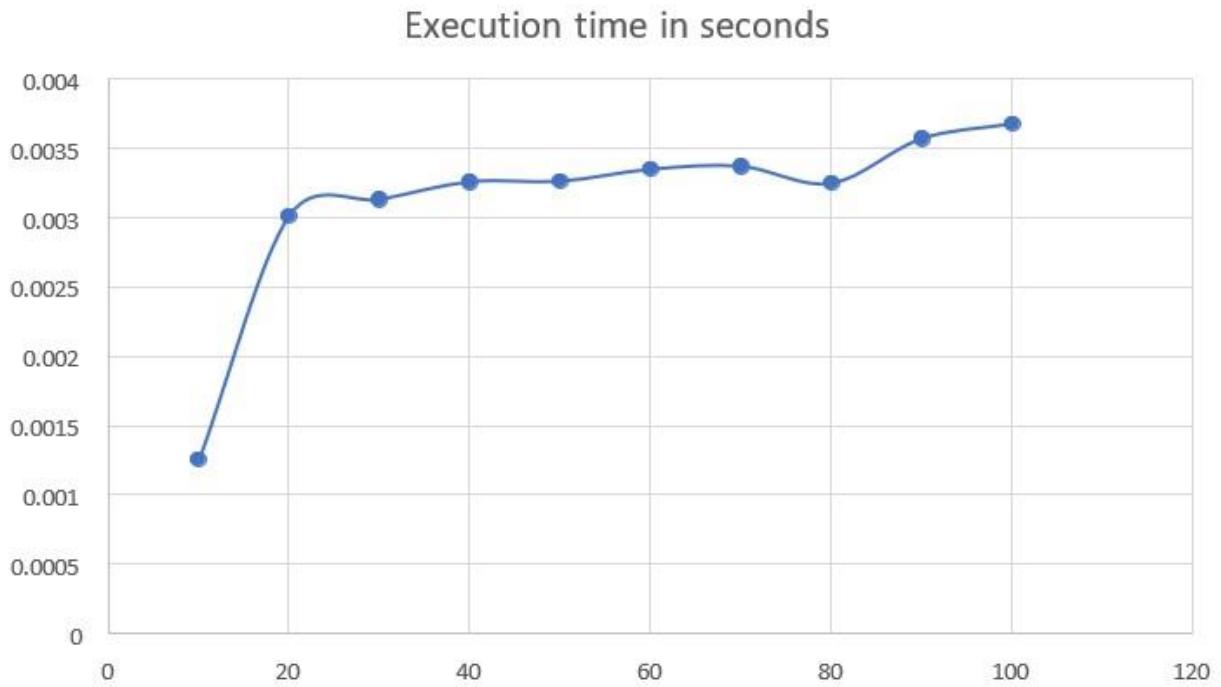
1 // C program to multiply two square matrices.
2 #include <stdio.h>
3 #include<math.h>
4 #include<omp.h>
5 #include<time.h>
6 #define N 4
7 //20BCE1988- O G RAGAVI
8 // This function multiplies mat1[][] and mat2[][][],
9 // and stores the result in res[][]
10
11 void multiply(int mat1[][N], int mat2[][N], int res[][N])
12 {
13     int i, j, k;
14     #pragma omp parallel for collapse(2) schedule(static)
15     for (i = 0; i < N; i++) {
16         for (j = 0; j < N; j++) {
17             res[i][j] = 0;
18             for (k = 0; k < N; k++)
19                 res[i][j] += mat1[i][k] * mat2[k][j];
20         }
21     }
22 }
23
24 int main()
25 {
26     clock_t s,e;
27     s=clock();
28     int mat1[N][N] = { { 1, 1, 1, 1 },
29                         { 2, 2, 2, 2 },
30                         { 3, 3, 3, 3 },
31                         { 4, 4, 4, 4 } };
32
33     int mat2[N][N] = { { 1, 1, 1, 1 },
34                         { 2, 2, 2, 2 },
35                         { 3, 3, 3, 3 },
36                         { 4, 4, 4, 4 } };
37
38     int res[N][N]; // To store result
39     int i, j;

```

## Output:

```
File Edit View Search Terminal Help  
ex2@AB1210SCOPE70:~$ gcc -o mul matrix_mul.c  
ex2@AB1210SCOPE70:~$ ./mul  
Result matrix is  
10 10 10 10  
20 20 20 20  
30 30 30 30  
40 40 40 40  
ex2@AB1210SCOPE70:~$ gcc -o mul matrix_mul.c  
ex2@AB1210SCOPE70:~$ ./mul  
Result matrix is  
10 10 10 10  
20 20 20 20  
30 30 30 30  
40 40 40 40  
ex2@AB1210SCOPE70:~$ gcc -o mul -fopenmp matrix_mul.c  
ex2@AB1210SCOPE70:~$ ./mul  
Result matrix is  
9 10 10 10  
20 20 20 20  
30 30 30 30  
40 60 40 40  
ex2@AB1210SCOPE70:~$ gcc -o mul -fopenmp matrix_mul.c  
ex2@AB1210SCOPE70:~$ ./mul  
Result matrix is  
10 10 10 10  
20 20 20 20  
30 30 30 30  
20 40 40 40  
Time taken 2726 ex2@AB1210SCOPE70:~$ gcc -o tsp tsp.c  
gcc: error: tsp.c: No such file or directory  
gcc: fatal error: no input files  
compilation terminated.  
ex2@AB1210SCOPE70:~$ gcc -o tsp tsp.c  
ex2@AB1210SCOPE70:~$ ./tsp.c  
bash: ./tsp.c: Permission denied  
ex2@AB1210SCOPE70:~$ ./tsp  
Enter Total Number of Cities: 4  
Enter Cost Matrix
```

## Graph:



## 2.OMP Program-Travelling salesman problem:

### Code:

```
#include <stdio.h>
#include<omp.h>
#include<time.h>
#include<stdlib.h>
//20BCE1988- O G RAGAVI
int matrix[25][25], visited_cities[10], limit, cost
= 0;

int tsp(int c)
{
int count, nearest_city =
999; int minimum = 999,
temp; for(count = 0; count
< limit; count++)
{ if((matrix[c][count] != 0) &&
(visited_cities[count] == 0))
{
```

```
if(matrix[c][count] < minimum)
{
    minimum = matrix[count][0] +
    matrix[c][count];
}
temp = matrix[c][count];
nearest_city = count;
}
}
if(minimum != 999)
{
    cost = cost + temp;
}
return nearest_city;
}
```

```
void minimum_cost(int city)
{ int
    nearest_city;
    visited_cities[cit
y] = 1;
    printf("%d ", city
+ 1);
    nearest_city =
    tsp(city);
    if(nearest_city
== 999)
    {
        nearest_city = 0; printf("%d",
        nearest_city + 1); cost =
        cost +
        matrix[city][nearest_city];
    }
}
```

```
}

minimum_cost(nearest_city);
}

int main()
{
int
i,
j;
clock_t s,e; s=clock();
printf("Enter Total Number of
Cities:\t"); scanf("%d",
&limit); printf("\nEnter Cost
Matrix\n");
for(i = 0; i < limit; i++)
{
printf("\nEnter %d Elements in Row[%d]\n",
limit, i + 1); for(j = 0; j < limit; j++)
{ scanf("%d",
&matrix[i][j]);
}
visited_cities
[i] = 0;
}

printf("\nEnterd Cost Matrix\n");
#pragma omp parallel for
for(i = 0; i < limit; i++)
{ printf("\n"); for(j
= 0; j < limit;
j++)
{ printf("%d ",
matrix[i][j]);
```

```

} } printf("\n\nPath:\t");
minimum_cost(0);
printf("\n\nMinimum Cost:
\t"); printf("%d\n", cost);
e=clock(); printf("\nTime
taken is %ld ",(e-s));
return 0; }

```

## Output:

```

File Edit View Search Terminal Help
ex2@AB1210SCOPE70:~$ gcc -o tsp -fopenmp tsp.c
tsp.c: In function `main':
tsp.c:81:2: warning: implicit declaration of function `prift'; did you mean `printf'? [-Wimplicit-function-declaration]
  81 |   prift("\nTime taken is %ld ",(e-s));
     |   ^
     |   printf
/usr/bin/ld: /tmp/ccPUTX0W.o: in function `main':
tsp.c:(.text+0x40f): undefined reference to `prift'
collect2: error: ld returned 1 exit status
ex2@AB1210SCOPE70:~$ gcc -o tsp -fopenmp tsp.c
ex2@AB1210SCOPE70:~$ ./tsp
Enter Total Number of Cities: 4

Enter Cost Matrix

Enter 4 Elements in Row[1]
1 2 3 4

Enter 4 Elements in Row[2]
5 6 7 8

Enter 4 Elements in Row[3]
3 4 5 6

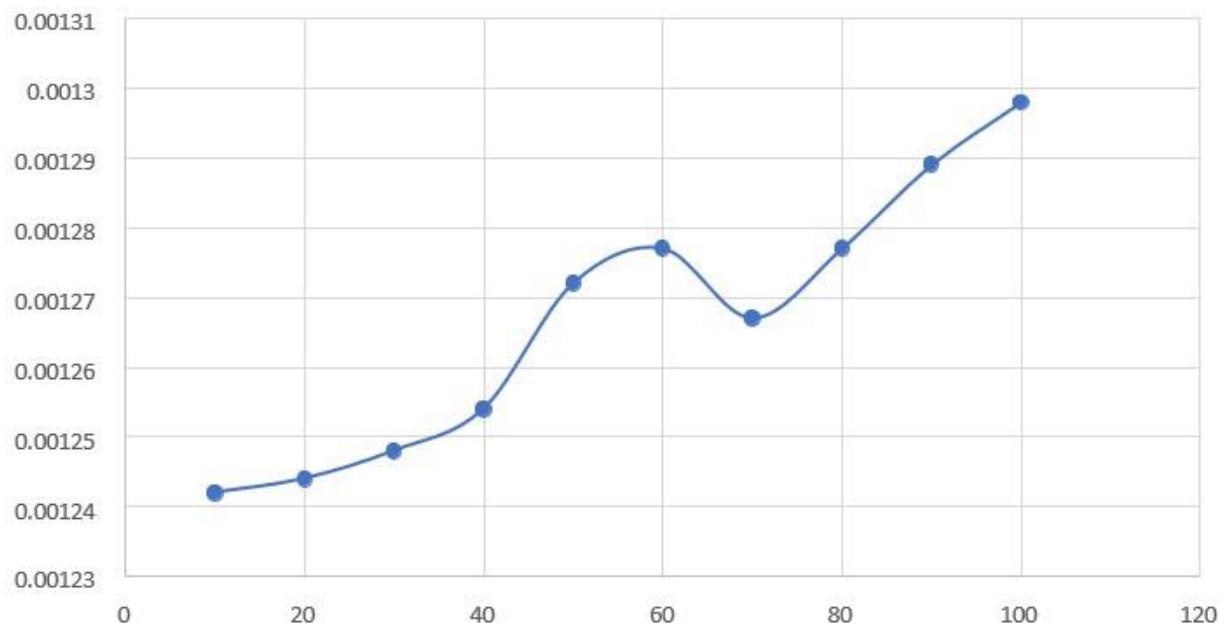
Enter 4 Elements in Row[4]
9 8 4 3

Entered Cost Matrix

3 4 5 6
1 2 3 4
9 8 4 3 5 6 7 8
Path: 1 4 3 2 1
Minimum Cost: 17
ex2@AB1210SCOPE70:~$ █

```

## Graph:



## **CSE-4001-Parallel and Distributed computing**

**Name:** O G RAGAVI

**Reg No:** 20BCE1988

**Title:** Scholarly articles on OpenMP

**Lab :** 7

---

---

**Scholarly article 1:** Research on OpenMP Model of the Parallel Programming Technology for Homogeneous Multicore DSP

### **Abstract:**

As application complexity continues to grow, using multicore processors has been proved to be an effective methodology to meet the ever-increasing processing demand across the industry association. The Master/Slave model, the Data Flow model and the OpenMP model are the three dominant models for parallel programming. In this paper, the first two models are briefly discussed while the OpenMP model is focused. Some factors (e.g. the number of threads, the scheduling strategy, the load balance, etc.) that affect the execution performance of OpenMP programs were also studied in this paper.

### **Synopsis:**

The use of parallel programming in the IT industry has grown significantly over the past few years.

Using the multi core processors to implement high performance computing has become the consensus of the industry. OpenMP is widely used in general purpose processors for its scalability and flexibility. Different hardware platforms correspond to different parallel programming models. The method using the OpenMP to realize the rapid and efficient image edge detection is presented in this paper. This paper presents a method of taking advantage of the OpenMP model to realize the image edge detection within the platform of TMS320C6678 Digital Signal Processors. The modification of the program only involves

the number of cores which proves the Open MP model has a good scalability.

## **Introduction:**

- Due to the limitation of the CMOS technology, it is difficult to continuously improve the frequency of the processor chip after it reaches 4GHz. Moore's Law leading the semiconductor market nearly 40 years may fail in the next 10 to 20 years. Using the multi core processors to implement high performance computing has become the consensus of the industry. Multicore Digital Signal Processors (DSP) are widely used in wireless telecommunication, industrial control, audio and video processing, etc. As embedded multicore hardware enables more functions to be implemented on the same device, efficient parallel programming methods are required to achieve desired performance without increasing software complexity.
- OpenMP is now widely used in general purpose processors for its scalability and flexibility, while the application in multicore DSP is still in the initial stage. The state-of-the-art approach to program multi-core DSPs is based on proprietary vendor Software Development Kits(SDKs), which only provides low-level, non-portable primitives
- Different hardware platforms correspond to different parallel programming models. This paper will take this platform as an instance, analyze the number of threads, the scheduling strategy on the performance of the OpenMP model. The method using the OpenMP to realize the rapid and efficient image edge detection is also presented in this paper.

## **Highlights:**

- This paper presents a method of taking advantage of the OpenMP model to realize the image edge detection within the platform of TMS320C6678 Digital Signal Processors (DSP)
- We study the two main factors which will influence the performance of  
OpenMP model

- The best performance can be obtained at the point where the number of threads is equal to the number of cores which are available within the platform
- In terms of image edge detection, we adopt the runtime scheduling strategy which gets a better performance compared to others
- The modification of the program only involves the number of cores which proves the OpenMP model has a good scalability.

## **Study subjects and analysis:**

The experimental results show that the OpenMP model has a better advantage on scalability and flexibility compared to the Master/Slave model and the Data Flow model. The Data Flow model is used for distributed control and execution **Results:**

TABLE I. EXECUTION TIME VERSUS NUMBER OF THREADS

	Number of threads		
	4	8	12
Execution Time/cycle	17112903	10353808	16902841

The table shows the text result that execution time varies with the number within the DSP platform. As can be seen, the execution time will bottom at a point where the number of threads is equal to the number of cores which are available within the platform. If the number of threads is less than the number of cores, then there will be some cores in idle state. Inversely, too many threads will result in the increasing overhead of threads creating and switching. Both conditions above are not conducive to performance improvement.

- When the SIZE is 80, then there are 80x80 multiply-add operations, the execution time is 19.2ms when the schedule(static) is specified. Since the main objective is to compare the pros and cons between various scheduling mechanisms, we want to calculate the relative value among them.

## Conclusion:

OpenMP is one of the most widely used parallel programming techniques in the modern multicore era. But applying this model to homogeneous multicore DSP is still a great challenge. In this paper, we studied the two main factors which will influence the performance of the OpenMP model.

The best performance can be obtained at the point where the number of threads is equal to the number of cores which are available within the platform. The scheduling strategy is determined by the specific application. In terms of image edge detection, we adopt the runtime scheduling strategy which gets a better performance compared to others. For the sake of validating the speedup, we use 1 to 8 cores to realize the edge detection and record the respective cost. The modification of the program only involves the number of cores which proves the OpenMP model has a good scalability. **The conclusion that the speedup increases linearly with the number of cores satisfies Gustafson's law. What's more, under the case of 8 cores running simultaneously, the speedup reaches 7.233.** Compared with the Master/Slave model, the OpenMP model has approximately improved one-third on performance.

## **Scholarly article2:** The Design of OpenMP Tasks

### **Abstract:**

OpenMP has been very successful in exploiting structured parallelism in applications. With increasing application complexity, there is a growing need for addressing irregular parallelism in the presence of complicated control structures. This is evident in various efforts by the industry and research communities to provide a solution to this challenging problem. One of the primary goals of OpenMP was to define a standard dialect to express and to exploit unstructured parallelism efficiently. This paper presents the design goals and key features of the tasking model, including a rich set of examples and an in-depth discussion of the rationale behind various design choices. It compares a prototype implementation of the tasking model with existing models, and evaluates it on a wide range of applications. The comparison shows that the OpenMP tasking model provides expressiveness, flexibility, and huge potential for performance and scalability

## **Introduction:**

In the last few decades, OpenMP has emerged as the de facto standard for shared-memory parallel programming. OpenMP provides a simple and flexible interface for developing portable and scalable parallel applications. OpenMP grew in the 1990s out of the need to standardize the different vendor specific directives related to parallelism. It was structured around parallel loops and was meant to handle dense numerical applications.

The set of features in the OpenMP 2.5 specification is ill equipped to exploit the concurrency available in modern applications. Users now need a simple way to identify independent units of work and not concern themselves with scheduling these work units. This model is typically called “tasking”. With the goal of defining a simple tasking dialect for expressing irregular and unstructured parallelism, a subcommittee of the OpenMP 3.0 language committee was formed.

## **Nested Parallelism Vs Tasking:**

New OpenMP tasks allow a programmer to express parallelism that in OpenMP 2.5 would be expressed using nested parallelism. The versions using nested OpenMP, while simple to write, usually do not perform well because of a variety of problems (load imbalance, synchronization overheads etc.). Handling recursive code structures is another simple case that motivated tasking in OpenMP was recursive work generation. Nested parallelism can be used to allow recursive work generation but at the expense of the overhead in creating a rigid tree structure of thread teams and their associated (unnecessary) implicit barriers.

## **Prototype implementation:**

In order to test the proposal in terms of expressiveness and performance, they have developed their own implementation of the proposed tasking model . They developed the prototype on top of a research OpenMP compiler (source-to-source restructuring tool) and runtime infrastructure. The runtime infrastructure is an implementation of a user level thread package based on the nano-threads programming model introduced first by Polychronopoulos.

The nano-thread layer is implemented on top of POSIX Threads (also known as pthreads). They decided to use pthreads to ensure that they will be portable across a wide range of systems. This layered implementation can have a slight impact on efficiency.

Once the task is first executed by a thread, and if the task has task scheduling points, we can expect two different behaviors. First, the task is bound to that thread (so, it can only be executed by that thread), and second, the task is not attached to any thread and can be executed by any other thread of the team. The library offers the possibility to move a task from the team queues to the local queues. This ability covers the requirements of the united clause of the task construct, which allows a task suspended by one thread to be resumed by a different one. The synchronization construct is provided through task counters that keep track of the number of tasks that are created in the current scope (i.e., the current task). Each task data structure has a successor field that points to the counter the task must decrement.

## **Conclusion:**

They have presented the work of the OpenMP 3.0 tasking subcommittee which was a proposal to **integrate task parallelism into the OpenMP specification**. This proposal allows programmers to **parallelize program structures like while loops and recursive functions more easily and efficiently**. We have shown that, in fact, these structures are easy to parallelize with the new proposal.

The process of defining the proposal has not been without difficult decisions, as they tried to achieve conflicting goals: simplicity of use, simplicity of specification, and consistency with the rest of OpenMP. Their discussions identified trade-offs between the goals, and our decisions reflected our best judgments of the relative merits of each. The comparisons of these results show that expressiveness is not incompatible with performance and the OpenMP tasks implementation can achieve very promising speedups when compared to other established models. Overall, **OpenMP tasks provide a balanced, flexible, and very expressive dialect for expressing unstructured parallelism** in OpenMP programs.

CSE4001-Parallel and  
distributed computing Name: OG RAGAVI  
Reg No: 20BCE1988  
Lab: 8

---

---

Steps:

## PARALLEL AND DISTRIBUTED COMPUTING

### Initial set up

```
sudo apt-get update
sudo apt-get install libopenmpi-dev
sudo apt-get install openmpi-bin

sudo apt-get install mpich
```

### Lab 8

#### 1. Hello World Program

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char** argv) {
    // Initialize the MPI environment
    MPI_Init(NULL, NULL);

    // Get the number of processes
    int world_size;
    MPI_Comm_size(MPI_COMM_WORLD, &world_size);

    // Get the rank of the process
    int world_rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);

    // Print off a hello world message
    printf("Hello world from rank %d out of %d processors\n",
        world_rank, world_size);
```

```
// Finalize the MPI environment.  
MPI_Finalize();  
}
```

**output:**

```
ragavi@ragavi-VirtualBox:~$ mpicc -o lab8_1 lab8_1.c  
ragavi@ragavi-VirtualBox:~$ mpirun -np 3 lab8_1  
Hello world from rank 0 out of 3 processors  
Hello world from rank 2 out of 3 processors  
Hello world from rank 1 out of 3 processors
```

## 2. Sum of Array Program

**CODE:**

```
#include <mpi.h>  
  
#include <stdio.h>  
  
#include <stdlib.h>  
  
#include <unistd.h>  
  
// size  
of  
array  
#defin  
e n 10  
  
int a[] = { 1, 2, 3, 4, 5, 6, 7,  
8, 9, 10 }; // Temporary  
array for slave process int  
a2[1000];  
  
int main(int argc, char* argv[])  
{  
    int pid, np,elements_per_process,n_elements_recieved;  
    // np -> no. of processes  
    // pid -> process id  
    MPI_Status status;
```

```

// Creation of parallel processes
MPI_Init(&argc, &argv);

// find out process ID,
// and how many processes were started
MPI_Comm_size(MPI_COMM_WORLD,&np);
MPI_Comm_rank(MPI_COMM_WORLD,&pid);

//  

m  

as  

te  

r  

pr  

oc  

es  

s  

if  

(p  

id  

=  

=  

0)
{  

int index, i;  

elements_per_process =n/np;  

// check if more than 1  

processes are run if (np > 1) {  

// distributes the portion of array  

// to child processes to calculate  

// their partial sums  

for (i = 1; i < np - 1;  

i++) { index = i *  

elements_per_proce  

ss;  

MPI_Send(&elements_per_process, 1, MPI_INT, i, 0,MPI_COMM_WORLD);  

MPI_Send(&a[index], elements_per_process,MPI_INT, i, 0,  

MPI_COMM_WORLD);

```

```

}

// last process adds remaining
elements index = i *
elements_per_process;

int elements_left = n - index;

MPI_Send(&elements_left, 1, MPI_INT,i, 0, MPI_COMM_WORLD);

MPI_Send(&a[index], elements_left,MPI_INT, i, 0, MPI_COMM_WORLD);

}

// master process add its own
sub array int sum = 0;

for (i = 0; i <
elements_per_process;i++) sum +=
a[i];

// collects partial sums from other
processes int tmp;

for (i = 1; i < np; i++) {

MPI_Recv(&tmp, 1, MPI_INT, MPI_ANY_SOURCE,
0,MPI_COMM_WORLD,
&status); int sender =
status.MPI_SOURCE; sum
+= tmp;

}

// prints the final sum of array
printf("Sum of array is : %d\n",
sum);

}

// slave
processes
else {

MPI_Recv(&n_elements_recieved, 1, MPI_INT, 0,0,MPI_COMM_WORLD,

```

```

    &status);

    // stores the received array segment
    // in local array a2

    MPI_Recv(&a2, n_elements_recieved, MPI_INT, 0, MPI_COMM_WORLD,
    &status);

    // calculates its
    partial sum int
    partial_sum = 0;

    for (int i = 0; i < n_elements_recieved; i++)
        partial_sum += a2[i];

    // sends the partial sum to the root process
    MPI_Send(&partial_sum, 1, MPI_INT, 0, MPI_COMM_WORLD);

}

// cleans up all MPI state before exit of process

MPI_Finalize();

return 0;
}

```

**output:**

```

ragavi@ragavi-VirtualBox:~$ mpicc -o lab8_2 lab8_2.c
ragavi@ragavi-VirtualBox:~$ mpirun -np 3 lab8_2
SUM OF array is : 55

```

## CSE4001-Parallel and distributed computing

**Name:** OG RAGAVI

**Reg No:** 20BCE1988

**Lab Ex:** 9

**Date:** 27.09.2022

---

### **Important commands:**

- **MPI\_Init (NULL, NULL)** –Initializes the MPI program.
- **MPI\_Comm\_size(MPI\_COMM\_WORLD,&no of processors)** – tells how many processors are there in communication.
- **int rank** – indicates the processor id according to its priority.
- **int MPI\_Send(void \*data\_to\_send, int send\_count, MPI\_Datatype send\_type, int destination\_ID, int tag, MPI\_Comm comm)**- used to send the data to the destination processor.Tag value depends on the number of different processes allocated to the system.If we have only one function, we'll use tag=0.
- **int MPI\_Recv(void \*received\_data, int receive\_count, MPI\_Datatype receive\_type,int sender\_ID, int tag, MPI\_Comm comm, MPI\_Status \*status)** –used to receive data from a particular processor.Status maintains the success result.

Q1) Program to sort an array of 100 numbers.

### **Algorithm:**

- Initialize the MPI program to run for multiple processors.
- Using rand() function , take input of 100 random numbers.
- Divide the program to np processors= no of elements in the array/np.

- Initially the master program gets into action and sorts the first 25 elements(in case of 100 numbers , as  $100/4$  processors= 25).
- Then it passes the index one by one (incrementing by 25) to the processor with pid 1, 2 and so on till np.
- Each slave computes the sorted array of their particular subpart. For eg)Pid=1 sorts the array from index 26 to 50 and so on.
- These partially sorted arrays are sent to the master processor , which are finally merged and sorted once again to display the output>

### **Code:**

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>
#include <mpi.h>

//sort for 100 numbers
void merge(int *, int *, int, int, int);
void mergeSort(int *, int *, int, int);

int main(int argc, char** argv) {

    /****** Create and populate the array *****/
    int n = atoi(argv[1]);
    int *original_array = malloc(n * sizeof(int));

    int c;
    srand(time(NULL));
}
```

```
printf("The initial unsorted array is: ");
for(c = 0; c < n; c++) {

    original_array[c] = rand() % n;
    printf("%d ", original_array[c]);

}

printf("\n");
printf("\n");

/********** Initialize MPI *****/
int world_rank;//priority given to the processor
int world_size;

MPI_Init(&argc, &argv);
MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);
MPI_Comm_size(MPI_COMM_WORLD, &world_size);

/********** Divide the array in equal-sized chunks *****/
int size = n/world_size;

/********** Send each subarray to each process *****/
int *sub_array = malloc(size * sizeof(int));
MPI_Scatter(original_array, size, MPI_INT, sub_array, size, MPI_INT, 0,
MPI_COMM_WORLD);
```

```
***** Perform the mergesort on each process *****/
int *tmp_array = malloc(size * sizeof(int));
mergeSort(sub_array, tmp_array, 0, (size - 1));

***** Gather the sorted subarrays into one *****/
int *sorted = NULL;
if(world_rank == 0) {

    sorted = malloc(n * sizeof(int));

}

MPI_Gather(sub_array, size, MPI_INT, sorted, size, MPI_INT, 0,
MPI_COMM_WORLD);

***** Make the final mergeSort call *****/
if(world_rank == 0) {

    int *other_array = malloc(n * sizeof(int));
    mergeSort(sorted, other_array, 0, (n - 1));

    ***** Display the sorted array *****/
    printf("The final sorted array is: ");
    for(c = 0; c < n; c++) {

        printf("%d ", sorted[c]);
    }
}
```

```
    }

    printf("\n");
    printf("\n");

    /****** Clean up root *****/
    free(sorted);
    free(other_array);

}

/****** Clean up rest *****/
free(original_array);
free(sub_array);
free(tmp_array);

/****** Finalize MPI *****/
MPI_Barrier(MPI_COMM_WORLD);
MPI_Finalize();

}

/****** Merge Function *****/
void merge(int *a, int *b, int l, int m, int r) {

    int h, i, j, k;
```

```
h = l;  
i = l;  
j = m + 1;  
  
while((h <= m) && (j <= r)) {  
  
    if(a[h] <= a[j]) {  
  
        b[i] = a[h];  
        h++;  
  
    }  
  
    else {  
  
        b[i] = a[j];  
        j++;  
  
    }  
  
    i++;  
  
}  
  
if(m < h) {
```

```
for(k = j; k <= r; k++) {  
  
    b[i] = a[k];  
    i++;  
  
}  
  
}  
  
else {  
  
    for(k = h; k <= m; k++) {  
  
        b[i] = a[k];  
        i++;  
  
    }
}  
  
}  
  
for(k = l; k <= r; k++) {  
  
    a[k] = b[k];  
  
}
}
```

```

}

/********** Recursive Merge Function *****/
void mergeSort(int *a, int *b, int l, int r) {

    int m;

    if(l < r) {

        m = (l + r)/2;

        mergeSort(a, b, l, m);
        mergeSort(a, b, (m + 1), r);
        merge(a, b, l, m, r);

    }

}

```

## **Output:**

```

ragevi@RAGAVI:~$ mpicc -o sort sort.c
ragevi@RAGAVI:~$ mpirun -np 1 ./sort 100
This is the unsorted array: 28 79 22 46 81 31 3 36 70 83 89 10 60 85 90 84 99 42 24 20 87 66 30 35 75 58 43 95 46 19 78 66 50 52 12 31 83 68 19 6 51 60 68 63 45 10 99 97 4 23 17 91 90 0 26 17 58 21 12 4
40 43 70 90 47 83 73 83 3 32 41 6 52 61 69 58 71 69 47 27 92 16 18 34 16 96 4 27 18 16 31 10 11 54 1 11 89 74 94 92

This is the sorted array: 0 1 3 3 4 4 4 6 6 10 10 11 11 12 12 16 16 16 17 17 18 18 19 19 20 20 21 22 23 24 26 27 27 30 31 31 34 35 36 40 41 42 43 43 45 46 46 47 47 50 50 51 52 52 54 58 58 60 60 61
63 66 66 68 68 69 69 70 70 71 73 74 75 75 78 79 81 83 83 83 84 85 87 89 89 90 90 91 92 92 94 95 96 97 99 99

```

## **Q2) Program to perform matrix multiplication:**

### **Code:**

```
#include <mpi.h>
#include <stdio.h>
#include<stdlib.h>
#include<math.h>
#define SIZE 8           /* Size of matrices */
```

```
int A[SIZE][SIZE], B[SIZE][SIZE], C[SIZE][SIZE];
```

```
void fill_matrix(int m[SIZE][SIZE])
{
    static int n=0;
    int i, j;
    for (i=0; i<SIZE; i++)
        for (j=0; j<SIZE; j++)
            //m[i][j] = n++;
            m[i][j]=rand()%50;
}
```

```
void print_matrix(int m[SIZE][SIZE])
{
    int i, j = 0;
    for (i=0; i<SIZE; i++) {
        printf("\n\t| ");
        for (j=0; j<SIZE; j++)
            printf("%2d ", m[i][j]);
        printf("|\n");
    }
}
```

```

}

}

int main(int argc, char *argv[])
{
    int myrank, P, from, to, i, j, k;
    int tag = 666;          /* any value will do */
    MPI_Status status;

    MPI_Init (&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &myrank);      /* who am i */
    MPI_Comm_size(MPI_COMM_WORLD, &P); /* number of processors */

    /* Just to use the simple variants of MPI_Gather and MPI_Scatter we */
    /* impose that SIZE is divisible by P. By using the vector versions, */
    /* (MPI_Gatherv and MPI_Scatterv) it is easy to drop this restriction. */

    if (SIZE%P!=0) {
        if (myrank==0) printf("Matrix size not divisible by number of processors\n");
        MPI_Finalize();
        exit(-1);
    }

    from = myrank * SIZE/P;
    to = (myrank+1) * SIZE/P;
}

```

```

/* Process 0 fills the input matrices and broadcasts them to the rest */
/* (actually, only the relevant stripe of A is sent to each process) */

if (myrank==0) {
    fill_matrix(A);
    fill_matrix(B);
}

MPI_Bcast (B, SIZE*SIZE, MPI_INT, 0, MPI_COMM_WORLD);

MPI_Scatter (A, SIZE*SIZE/P, MPI_INT, A[from], SIZE*SIZE/P, MPI_INT, 0,
MPI_COMM_WORLD);

printf("Result of matrix multiplication is:");
for (i=from; i<to; i++)
    for (j=0; j<SIZE; j++) {
        C[i][j]=0;
        for (k=0; k<SIZE; k++)
            C[i][j] += A[i][k]*B[k][j];
    }

MPI_Gather (C[from], SIZE*SIZE/P, MPI_INT, C, SIZE*SIZE/P, MPI_INT, 0,
MPI_COMM_WORLD);

if (myrank==0) {
    printf("\n\n");
    print_matrix(A);
}

```

```

printf("\n\n\t * \n");
print_matrix(B);
printf("\n\n\t = \n");
print_matrix(C);
printf("\n\n");
}


```

```

MPI_Finalize();
return 0;
}
```

## **Output:**

ragavi@RAGAVI:~\$ mpicc -o mul1 mat\_mul.c  
 ragavi@RAGAVI:~\$ mpirun -np 1 mul1  
 Result of matrix multiplication is:

33 36 27 15 43 35 36 42
49 21 12 27 40 9 13 26
46 26 22 36 11 18 17 29
32 30 12 23 17 35 29 2
22 8 19 17 43 6 11 42
29 23 21 19 34 37 48 24
15 20 13 26 41 30 6 23
12 20 46 31 5 25 34 27

\*

36 5 46 29 13 7 24 45
32 45 14 17 34 14 43 0
37 8 26 28 38 34 3 1
4 49 32 10 26 18 39 12
26 36 44 39 45 20 34 28
17 1 47 2 17 42 2 6
1 30 36 41 15 39 44 19
40 29 31 17 47 21 31 25

=

6828 6617 9339 6412 8113 6539 7424 4840
5234 5202 7181 4937 5879 3762 5924 4612
4999 5075 6769 4233 5595 4288 5605 3718
3794 4308 6439 3960 4379 4449 4977 3015
4730 4557 6034 4330 5898 3641 4852 3712
5154 5676 8517 5786 6458 6191 6575 4240
4267 4706 6283 3695 5571 4092 4704 3017
4567 4855 6476 4384 5819 5529 5048 2569

ragavi@RAGAVI:~\$

## CSE4001-Parallel and distributed computing

Name: O G RAGAVI

Reg No: 20BCE1988

Lab Ex: 10

Title: Dijshktra's algorithm implementation using open MP and mpi

---

### Open MP implementation:

#### Code:

```
#include<omp.h>
#include <limits.h>
#include <stdbool.h>
#include <stdio.h>
#include<time.h>

// Number of vertices in the graph
#define V 9

int minDistance(int dist[], bool sptSet[])
{
    // Initialize min value    int min =
    INT_MAX, min_index;

    for (int v = 0; v < V; v++) if (sptSet[v] == false && dist[v] <=
        min)          #pragma omp critical           min
        = dist[v], min_index = v;

    return min_index;
}

// A utility function to print the constructed distance
// array
void printSolution(int dist[])
{
}
```

```

printf("Vertex \t\t Distance from Source\n");    for (int i = 0;
i < V; i++)           printf("%d \t\t\t %d\n", i, dist[i]);
}

void dijkstra(int graph[V][V], int src)
{
    int dist[V];

    bool sptSet[V];
    for (int i = 0; i < V; i++)
        dist[i] = INT_MAX, sptSet[i] = false;

    dist[src] = 0;

    #pragma omp parallel for

    for (int count = 0; count < V - 1; count++) {

        int u = minDistance(dist, sptSet);

        // Mark the picked vertex as processed      sptSet[u] = true;

        for (int v = 0; v < V; v++)

            if (!sptSet[v] && graph[u][v]
                && dist[u] != INT_MAX
                && dist[u] + graph[u][v] < dist[v])

```

```

        dist[v] = dist[u] + graph[u][v];

    }

    printSolution(dist);

}

// driver's code int
main()
{
    int graph[V][V] = { { 0, 4, 0, 0, 0, 0, 0, 8, 0 },
                        { 4, 0, 8, 0, 0, 0, 0, 11, 0 },
                        { 0, 8, 0, 7, 0, 4, 0, 0, 2 },
                        { 0, 0, 7, 0, 9, 14, 0, 0, 0 },
                        { 0, 0, 0, 9, 0, 10, 0, 0, 0 },
                        { 0, 0, 4, 14, 10, 0, 2, 0, 0 },
                        { 0, 0, 0, 0, 0, 2, 0, 1, 6 },
                        { 8, 11, 0, 0, 0, 0, 1, 0, 7 },
                        { 0, 0, 2, 0, 0, 0, 6, 7, 0 } };

    // Function call
    clock_t s,e;
    s=clock();
    dijkstra(graph, 0);      e=clock();
    printf("\nfor a 9x9 adjacency matrix:");
    printf("\nTime taken %ld\n", (e-s));

    return 0;
}

```

**Output:**

```
ragavi@RAGAVI:~/Documents$ ./a.out
Vertex          Distance from Source
0                  0
1                  4
2                 12
3                 19
4                 21
5                 11
6                  9
7                  8
8                 14

for a 9x9 adjacency matrix:
Time taken 527
```

```
ragavi@RAGAVI:~/Documents$ gcc -fopenmp lab10a.c
ragavi@RAGAVI:~/Documents$ ./a.out
Vertex          Distance from Source
0                  0
1                  1
2                  1
3                  1

for a 4x4 adjacency matrix:
Time taken 50
ragavi@RAGAVI:~/Documents$
```

**2.MPI Implementation:**

**Code:**

```
#include <stdlib.h>
#include <stdio.h>
#include <string.h>
#include <math.h>
#include <sys/time.h>
#include "mpi.h"

#define N 5
#define SOURCE 1
#define MAXINT 9999999

void SingleSource(int n, int source, int *wgt, int *lengths, MPI_Comm comm)
{
    int temp[N];
```

```

int i, j;

int nlocal; /* The number of vertices stored locally */ int *marker; /* Used to mark
the vertices belonging to Vo */

int firstvtx; /* The index number of the first vertex that is stored locally
*/
int lastvtx; /* The index number of the last vertex that is stored locally
*/ int u,
udist; int
lminpair[2],
gminpair[2];
int npes,
myrank;
MPI_Status status;

MPI_Comm_size(comm, &npes); MPI_Comm_rank(comm, &myrank);
nlocal = n / npes; firstvtx = myrank *
nlocal; lastvtx = firstvtx + nlocal - 1;
/* Set the initial distances from source to all the other vertices */ for (j = 0; j < nlocal; j++)
{
lengths[j] = wgt[source * nlocal + j];
}

/* This array is used to indicate if the shortest part to a vertex has been found or not.*/
/* if marker [v] is one, then the shortest path to v has been found, otherwise it has not.*/
marker = (int *)malloc(nlocal * sizeof(int)); for (j = 0; j < nlocal;
j++)
{
marker[j] = 1;
}

/* The process that stores the source vertex, marks it as being seen */ if (source >= firstvtx &&
source <= lastvtx)
{
marker[source - firstvtx] = 0;
}

/* The main loop of Dijkstra's algorithm */ for (i = 1; i < n; i++)

```

```

{
/* Step 1: Find the local vertex that is at the smallest distance from source */
lminpair[0] = MAXINT; /* set it to an architecture dependent large number
*/
lminpair[1] = -1; for (j = 0; j <
nlocal; j++)
{
if (marker[j] && lengths[j] < lminpair[0])
{
lminpair[0] = lengths[j]; lminpair[1] = firstvtx +
j;
}
}
/* Step 2: Compute the global minimum vertex, and insert it into Vc */
MPI_Allreduce(lminpair, gminpair, 1, MPI_2INT, MPI_MINLOC, comm); udist = gminpair[0]; u =
gminpair[1];
/* The process that stores the minimum vertex, marks it as being seen
*/
if (u == lminpair[1])
{
marker[u - firstvtx] = 0;
}
/* Step 3: Update the distances given that u got inserted */ for (j = 0; j < nlocal; j++)
{
if (marker[j] && ((udist + wgt[u * nlocal + j]) < lengths[j]))
{
lengths[j] = udist + wgt[u * nlocal + j];
}
}
}
}
free(marker);
}

```

```

int main(int argc, char *argv[])
{
    int npes, myrank, nlocal; int weight[N][N]; /*adjacency matrix*/
    int distance[N]; /*distance vector*/ int *localWeight; /*local weight
array*/ int *localDistance; /*local distance vector*/ int sendbuf[N
* N]; /*local weight to distribute*/ int i, j, k;
    char fn[255];

    double time_start, time_end; struct
    timeval tv; struct timezone tz;
    gettimeofday(&tv, &tz);
    time_start = (double)tv.tv_sec + (double)tv.tv_usec / 1000000.00;

    /* Initialize MPI and get system information */
    MPI_Init(&argc, &argv);

    MPI_Comm_size(MPI_COMM_WORLD, &npes);

    MPI_Comm_rank(MPI_COMM_WORLD, &myrank);

    nlocal = N / npes; /* Compute the number of elements to be stored locally. */

    /*allocate local weight and local disatance arrays for each prosess*/ localWeight = (int
    *)malloc(nlocal * N * sizeof(int));

    localDistance = (int *)malloc(nlocal * sizeof(int));

    /* Open input file, read adjacency matrix and prepare for sendbuf */

    // printf("\nThe adjacency matrix: \n"); for (i = 0; i < N;
    i++)
    {
        for (j = 0; j < N; j++)
        {
            weight[i][j] = rand() % 11;
            // if (weight[i][j] == 9999999) printf("%4s", "INT");
            // else printf("%4d", weight[i][j]);
        }
        // printf("\n");
    }
}

```

```

/*prepare send data */ for (k = 0; k
< npes; ++k)
{
for (i = 0; i < N; ++i)
{
for (j = 0; j < nlocal; ++j)
{
sendbuf[k * N * nlocal + i * nlocal + j] = weight[i][k * nlocal + j];
}
}
}

/*distribute data*/
MPI_Scatter(sendbuf, nlocal * N, MPI_INT, localWeight, nlocal * N,
MPI_INT, SOURCE,
MPI_COMM_WORLD);

/*Implement the single source dijkstra's algorithm*/
SingleSource(N, SOURCE, localWeight, localDistance,
MPI_COMM_WORLD);

/*collect local distance vector at the source process*/
MPI_Gather(localDistance, nlocal, MPI_INT, distance, nlocal, MPI_INT,
SOURCE,
MPI_COMM_WORLD); if
(myrank == SOURCE)
{
printf("Nodes: %d\n", N); printf("The distance
vector is \n"); for (i = 0; i < N; ++i)
{
printf("%d ", distance[i]);
}
printf("\n");
gettimeofday(&tv, &tz);
time_end = (double)tv.tv_sec + (double)tv.tv_usec / 1000000.00;

```

```
printf("time cost is %1f\n", time_end - time_start);  
}  
  
free(localWeight);  
free(localDistance);  
  
MPI_Finalize(); return 0;  
}
```

**Output:**

```
ragavi@RAGAVI:~$ mpicc -o x x.c  
ragavi@RAGAVI:~$ mpirun -np 2 ./x  
Nodes: 5  
The distance vector is  
4 0 6 3 0  
time cost is 0.230125
```

# CSE4001-Parallel and distributed computing

**Name:** OG RAGAVI

**Reg No:** 20BCE1988

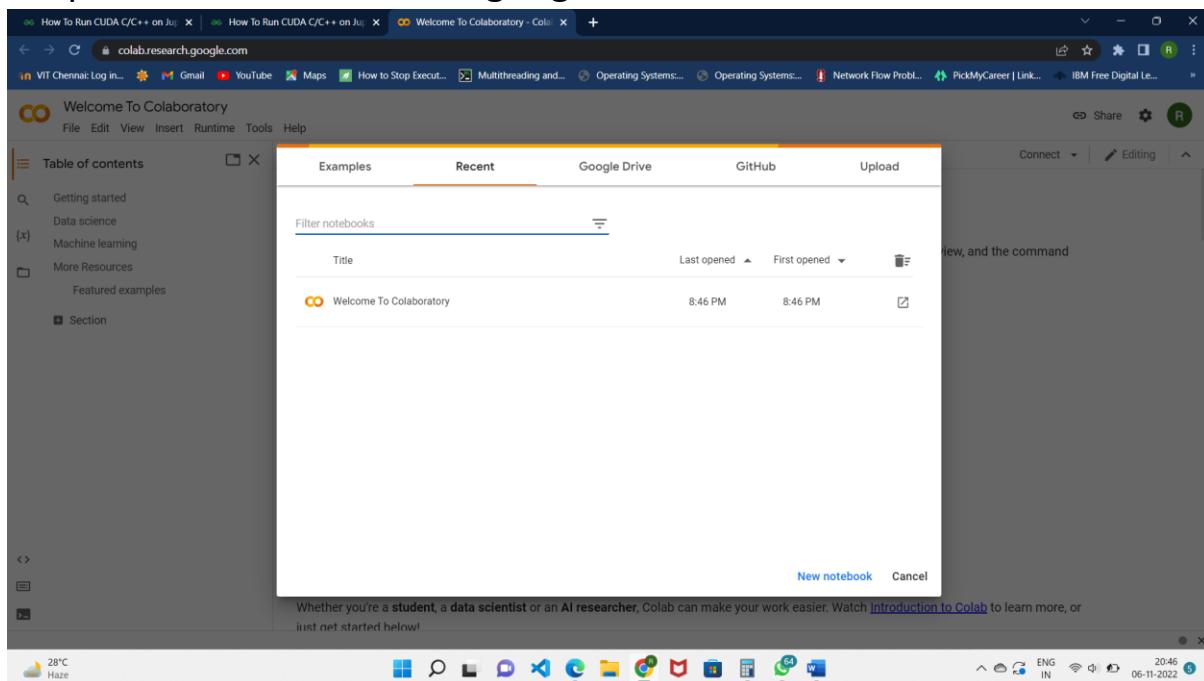
**Lab Ex:** 11

**Title:** CUDA C/C++

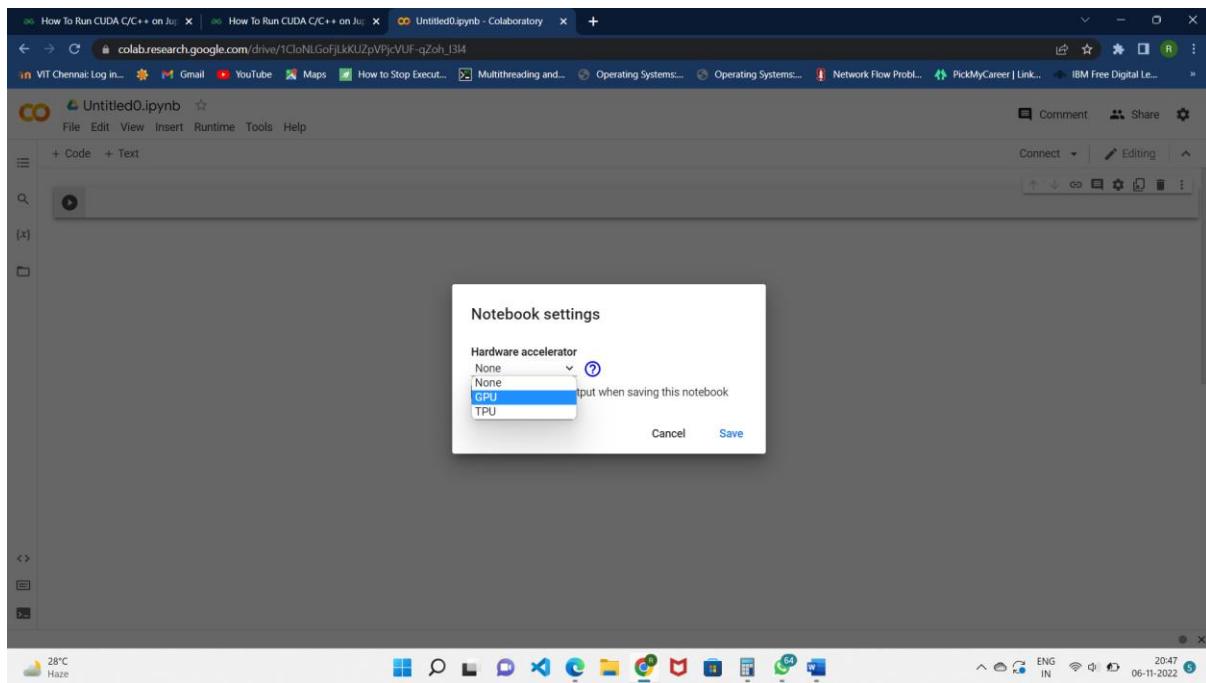
---

## Installation of CUDA C/C++ Google Colab:

1. Open a new notebook in google colab



## 2.Switch runtime to GPU:



## 3.To uninstall previous CUDA version:

```
!apt-get --purge remove cuda nvidia* libnvidia-*
```

```
!dpkg -l | grep cuda- | awk '{print $2}' | xargs -n1 dpkg --purge
```

```
!apt-get remove cuda-*
```

```
!apt autoremove
```

```
!apt-get update
```

How To Run CUDA C/C++ on Jupyter Notebook

RAGAVI\_20BCE1988.ipynb - Colab

```
+ Code + Text
apt-get --purge remove cuda nvidia* libnvidia-*  

keyboard-configuration libargon2-0 libcap2 libcryptsetup2  

libdevmapper1.02.1 libfontenc1 libidn11 libip4tc0 libjansson libpam-systemd  

libpolkit-agent-1-0 libpolkit-backend-1-0 libpolkit-gobject-1-0 libxfont2  

libxkbfile1 libxvtr10 libxtst6 nisight-compute-2020.3.1  

nisight-systems-2020.4.3 ocl-icd-opencl-dev opencl-c-headers openjdk-11-jre  

policykit-1 policykit-1-gnome python3-xkit screen-resolution-extra systemd  

systemd-sysv udev x11-xkb-utils xserver-common xserver-xorg-core-hwe-18.04  

Use 'apt autoremove' to remove them.  

The following packages will be REMOVED:  

cuda-11-2* cuda-demo-suite-11-2* cuda-drivers* cuda-drivers-520*  

cuda-runtime-11-2* libnvidia-cfg1-520* libnvidia-common-460*  

libnvidia-common-520* libnvidia-compute-460* libnvidia-compute-520*  

libnvidia-decode-520* libnvidia-encode-520* libnvidia-extra-520*  

libnvidia-fbc1-520* libnvidia-gl-520* nvidia-compute-utils-520*  

nvidia-dkms-520* nvidia-driver-520* nvidia-kernel-common-520*  

nvidia-kernel-source-520* nvidia-modprobe* nvidia-opencl-dev*  

nvidia-settings* nvidia-utils-520* xserver-xorg-video-nvidia-520*  

0 upgraded, 0 newly installed, 25 to remove and 4 not upgraded.  

After this operation, 824 MB disk space will be freed.  

(Reading database ... 123942 files and directories currently installed.)  

Removing cuda-11-2 (11.2.2-1) ...  

Removing cuda-demo-suite-11-2 (11.2.152-1) ...  

Removing cuda-runtime-11-2 (11.2.2-1) ...  

Removing cuda-drivers* (520.61.05-1) ...  

Removing cuda-drivers-520 (520.61.05-1) ...  

Removing nvidia-driver-520 (520.61.05-ubuntu1) ...  

Removing xserver-xorg-video-nvidia-520 (520.61.05-ubuntu1) ...  

Removing libnvidia-cfg1-520:amd64 (520.61.05-ubuntu1) ...  

Removing libnvidia-common-460 (460.106.00-ubuntu1) ...  

Removing libnvidia-encode-520:amd64 (520.61.05-ubuntu1) ...  

Removing libnvidia-gl-520:amd64 (520.61.05-ubuntu1) ...  

Removing libnvidia-common-520 (520.61.05-ubuntu1) ...  

Removing libnvidia-utils-520 (520.61.05-ubuntu1) ...  

Removing nvidia-settings* (520.61.05-ubuntu1) ...  

dpkg: dependency problems prevent removal of cuda-command-line-tools-11-2:  

  cuda-tools-11-2 depends on cuda-command-line-tools-11-2 (>= 11.2.2).  

dpkg: error processing package cuda-command-line-tools-11-2 (--purge):  

  dependency problems - not removing  

Errors were encountered while processing:  

  cuda-command-line-tools-11-2  

(Reading database ... 123231 files and directories currently installed.)  

Removing cuda-compat-11-2 (460.106.00-1) ...  

Processing triggers for libc-bin (2.27-3ubuntu1.6) ...  

dpkg: dependency problems prevent removal of cuda-compiler-11-2:  

  cuda-toolkit-11-2 depends on cuda-compiler-11-2 (>= 11.2.2).  

  cuda-minimal-build-11-2 depends on cuda-compiler-11-2 (>= 11.2.2).  

dpkg: error processing package cuda-compiler-11-2 (--purge):  

  dependency problems - not removing  

Errors were encountered while processing:  

  cuda-compiler-11-2  

dpkg: dependency problems prevent removal of cuda-cudart-11-2:  

  cuda-cudart-dev-11-2 depends on cuda-cudart-11-2 (>= 11.2.152).  

  cuda-libraries-11-2 depends on cuda-cudart-11-2 (>= 11.2.152).  

dpkg: error processing package cuda-cudart-11-2 (--purge):  

  dependency problems - not removing  

Errors were encountered while processing:  

  cuda-cudart-11-2  

dpkg: dependency problems prevent removal of cuda-cudart-dev-11-2:  

  cuda-nvcc-11-2 depends on cuda-cudart-dev-11-2.  

  cuda-libraries-dev-11-2 depends on cuda-cudart-dev-11-2 (>= 11.2.152).  

  cuda-cupti-11-2 depends on cuda-cudart-dev-11-2.  

  cuda-visual-tools-11-2 depends on cuda-cudart-dev-11-2 (>= 11.2.152).  

2s completed at 8:51 PM
```

27°C Satisfactory air

ENG IN 2051 06-11-2022

How To Run CUDA C/C++ on Jupyter Notebook

RAGAVI\_20BCE1988.ipynb - Colab

```
+ Code + Text
dpkg -1 | grep cuda- | awk '{print $2}' | xargs -n1 dpkg --purge  

dpkg: dependency problems prevent removal of cuda-command-line-tools-11-2:  

  cuda-tools-11-2 depends on cuda-command-line-tools-11-2 (>= 11.2.2).  

dpkg: error processing package cuda-command-line-tools-11-2 (--purge):  

  dependency problems - not removing  

Errors were encountered while processing:  

  cuda-command-line-tools-11-2  

(Reading database ... 123231 files and directories currently installed.)  

Removing cuda-compat-11-2 (460.106.00-1) ...  

Processing triggers for libc-bin (2.27-3ubuntu1.6) ...  

dpkg: dependency problems prevent removal of cuda-compiler-11-2:  

  cuda-toolkit-11-2 depends on cuda-compiler-11-2 (>= 11.2.2).  

  cuda-minimal-build-11-2 depends on cuda-compiler-11-2 (>= 11.2.2).  

dpkg: error processing package cuda-compiler-11-2 (--purge):  

  dependency problems - not removing  

Errors were encountered while processing:  

  cuda-compiler-11-2  

dpkg: dependency problems prevent removal of cuda-cudart-11-2:  

  cuda-cudart-dev-11-2 depends on cuda-cudart-11-2 (>= 11.2.152).  

  cuda-libraries-11-2 depends on cuda-cudart-11-2 (>= 11.2.152).  

dpkg: error processing package cuda-cudart-11-2 (--purge):  

  dependency problems - not removing  

Errors were encountered while processing:  

  cuda-cudart-11-2  

dpkg: dependency problems prevent removal of cuda-cudart-dev-11-2:  

  cuda-nvcc-11-2 depends on cuda-cudart-dev-11-2.  

  cuda-libraries-dev-11-2 depends on cuda-cudart-dev-11-2 (>= 11.2.152).  

  cuda-cupti-11-2 depends on cuda-cudart-dev-11-2.  

  cuda-visual-tools-11-2 depends on cuda-cudart-dev-11-2 (>= 11.2.152).  

2s completed at 8:51 PM
```

27°C Satisfactory air

ENG IN 2052 06-11-2022

How To Run CUDA C/C++ on Jupyter Notebook | Colab

File Edit View Insert Runtime Tools Help All changes saved

```
+ Code + Text
Q lapt-get remove cuda-*
```

[x] Package 'cuda-cudart-11-0' is not installed, so not removed  
Package 'cuda-drivers-fabricmanager-515' is not installed, so not removed  
Package 'cuda-drivers-fabricmanager-450' is not installed, so not removed  
Package 'cuda-drivers-fabricmanager-460' is not installed, so not removed  
Package 'cuda-drivers-fabricmanager-470' is not installed, so not removed  
Package 'cuda-drivers-fabricmanager-510' is not installed, so not removed  
The following packages were automatically installed and are no longer required:  
default-jre dkms keyboard-configuration libargon2-0 libcapt2 libcryptsetup12  
libcuftt-11-2 libcuftt-dev-11-2 libcurand-11-2 libcurand-dev-11-2  
libcusolver-11-2 libcusolver-dev-11-2 libdevmapper1.02.1 libfontenc1  
libidn11 libip4t0 libjansson libnvjpeg-11-2 libnvjpeg-dev-11-2  
libpam-system libpolkit-agent-1-0 libpolkit-backend-1-0  
libpolkit-gobject-1-0 libxfont2 libxbkfst1 libxnvctrl libxtst6  
nsight-compute-2020.3.1 nsight-systems-2020.4.3 ocl-icd-opencl-dev  
openrc-helpers openrc-udev policykit-1 policykit-1-gnome python3-xkit  
screen-resolution-extra systemd systemd-sysv udev x11-xkb-utils  
xserver-common xserver-xorg-core-18.04  
Use 'apt autoremove' to remove them.  
The following packages will be REMOVED:  
cuda-command-line-tools-11-2 cuda-compiler-11-2 cuda-cudart-11-2  
cuda-cudart-dev-11-2 cuda-cubjdump-11-2 cuda-cupti-11-2 cuda-cupti-dev-11-2  
cuda-cuxfilt-11-2 cuda-documentation-11-2 cuda-driver-dev-11-2  
cuda-gdb-11-2 cuda-libraries-11-2 cuda-libraries-dev-11-2 cuda-memcheck-11-2  
cuda-nsight-11-2 cuda-nsight-compute-11-2 cuda-nsight-systems-11-2  
cuda-nvcc-11-2 cuda-nvidasm-11-2 cuda-nvml-dev-11-2 cuda-nvprof-11-2  
cuda-nvprune-11-2 cuda-nvrtc-11-2 cuda-nvrtc-dev-11-2 cuda-nvtx-11-2  
cuda-nvvp-11-2 cuda-samples-11-2 cuda-sanitizer-11-2  
@ upgraded, @ newly installed, 28 to remove and @ not upgraded.  
After this operation, 938 MB disk space will be freed.  
(Reading database ... 123209 files and directories currently installed.)

5s completed at 8:52 PM

27°C AQI 91.96 ENG IN 2052 06-11-2022

How To Run CUDA C/C++ on Jupyter Notebook | Colab

File Edit View Insert Runtime Tools Help Saving...

```
+ Code + Text
Q lapt autoremove
```

[x] screen-resolution-extra systemd-sysv udev x11-xkb-utils  
xserver-common xserver-xorg-core-hwe-18.04  
@ upgraded, @ newly installed, 43 to remove and @ not upgraded.  
After this operation, 2,237 kB disk space will be freed.  
(Reading database ... 117860 files and directories currently installed.)  
Removing default-jre (2:1.11-6ubuntu18.04.1) ...  
Removing dkms (2.3-3ubuntu9.7) ...  
Removing xserver-xorg-core-hwe-18.04 (2:1.20.8-8ubuntu2.2-18.04.7) ...  
Removing keyboard-configuration (1.17ubuntu2.9) ...  
Removing screen-resolution-extra (0.17.3) ...  
Removing policykit-1-gnome (0.105-6ubuntu2) ...  
Removing policykit-1 (0.105-20ubuntu0.18.04.6) ...  
Removing libpam-systemd:amd64 (237-3ubuntu10.56) ...  
Removing systemd-sysv (237-3ubuntu10.56) ...  
Removing system (237-3ubuntu10.56) ...  
Removing libcryptsetup12:amd64 (2:2.0.2-1ubuntu1.2) ...  
Removing libargon2-0:amd64 (0-20161029-1.1) ...  
Removing libcap2:amd64 (1:2.25-1.2) ...  
Removing libcuftt-dev-11-2 (10.4.1.152-1) ...  
Removing libcuftt-11-2 (10.4.1.152-1) ...  
Removing libcurand-dev-11-2 (10.2.3.152-1) ...  
Removing libcurand-11-2 (10.2.3.152-1) ...  
Removing libcusolver-dev-11-2 (11.1.0.152-1) ...  
Removing libcusolver-11-1 (11.1.0.152-1) ...  
Removing libdevmapper1.02.1:amd64 (2:1.02.145-4.1ubuntu3.18.04.3) ...  
Removing libfont2:amd64 (1:2.0.3-1) ...  
Removing libfontenc1:amd64 (1:1.1.3-1) ...  
Removing libidn11:amd64 (1.33-2.1ubuntu1.2) ...  
Removing libip4t0:amd64 (1.6.1-2ubuntu2) ...  
Removing libjansson4:amd64 (2.11-1) ...  
Removing libnvjpeg-dev-11-2 (11.4.0.152-1) ...

6s completed at 8:53 PM

27°C AQI 91.96 ENG IN 2053 06-11-2022

```
!apt-get update
```

```
[x] Hit:1 http://archive.ubuntu.com/ubuntu bionic InRelease  
Get:2 http://archive.ubuntu.com/ubuntu bionic-updates InRelease [88.7 kB]  
Get:3 https://cloud.r-project.org/bin/linux/ubuntu bionic-cran40/ InRelease [3,626 B]  
Get:4 http://archive.ubuntu.com/ubuntu bionic-backports InRelease [83.3 kB]  
Get:5 http://security.ubuntu.com/ubuntu bionic-security InRelease [88.7 kB]  
Get:6 https://ppa.launchpad.net/c2d4u-team/c2d4u_0.0/+ubuntu bionic InRelease [15.9 kB]  
Ign:7 https://developer.download.nvidia.com/compute/machine-learning/repos/ubuntu1804/x86_64 InRelease  
Hit:8 https://developer.download.nvidia.com/compute/machine-learning/repos/ubuntu1804/x86_64 Release  
Hit:9 https://ppa.launchpad.net/cran/l1hpbit/ubuntu bionic InRelease  
Get:10 https://cloud.r-project.org/bin/linux/ubuntu bionic-cran40/ Packages [98.9 kB]  
Get:11 https://ppa.launchpad.net/deadsnakes/ppa/ubuntu bionic InRelease [15.9 kB]  
Get:12 http://archive.ubuntu.com/ubuntu bionic-updates/universe amd64 Packages [2,332 kB]  
Get:13 http://archive.ubuntu.com/ubuntu bionic-updates/main amd64 Packages [3,472 kB]  
Hit:14 http://ppa.launchpad.net/graphics-drivers/ppa/ubuntu bionic InRelease  
Get:15 http://archive.ubuntu.com/ubuntu bionic-updates/restricted amd64 Packages [1,267 kB]  
Get:16 http://ppa.launchpad.net/c2d4u-team/c2d4u_0.0/+ubuntu bionic/main Sources [2,216 kB]  
Get:18 http://security.ubuntu.com/ubuntu bionic-security/main amd64 Packages [3,040 kB]  
Get:19 http://ppa.launchpad.net/c2d4u-team/c2d4u_0.0/+ubuntu bionic/main amd64 Packages [1,133 kB]  
Get:20 http://security.ubuntu.com/ubuntu bionic-security/restricted amd64 Packages [1,226 kB]  
Get:21 http://security.ubuntu.com/ubuntu bionic-security/universe amd64 Packages [1,554 kB]  
Fetched 16.6 MB in 3s (6,362 kB/s)  
Reading package lists... Done
```

To install CUDA version 9:

**!wget**

**https://developer.nvidia.com/compute/cuda/9.2/Prod/local\_installers/cuda-repo-ubuntu1604-9-2-local\_9.2.88-1\_amd64 -O cuda-repo-ubuntu1604-9-2-local\_9.2.88-1\_amd64.deb**

**!dpkg -i cuda-repo-ubuntu1604-9-2-local\_9.2.88-1\_amd64.deb**

**!apt-key add /var/cuda-repo-9-2-local/7fa2af80.pub**

**!apt-get update**

**!apt-get install cuda-9.2**

```

!wget https://developer.nvidia.com/compute/cuda/9.2/Prod/local_installers/cuda-repo-ubuntu1604-9-2-local_9.2.88-1_amd64 -O cuda-repo-ubuntu1604-9-2-local_9.2.88-1_amd64.deb
!dpkg -i cuda-repo-ubuntu1604-9-2-local_9.2.88-1_amd64.deb
!apt-key add /var/cuda-repo-9-2-local/7faaf80.pub
!apt-get update
!apt-get install cuda-9.2

--2022-11-06 15:24:57- https://developer.nvidia.com/compute/cuda/9.2/Prod/local_installers/cuda-repo-ubuntu1604-9-2-local_9.2.88-1_amd64
Resolving developer.nvidia.com (developer.nvidia.com)... 152.195.19.142
Connecting to developer.nvidia.com (developer.nvidia.com)|152.195.19.142|:443... connected.
HTTP request sent, awaiting response... 301 Moved Permanently
Location: https://developer.download.nvidia.com/compute/cuda/9.2/secure/Prod/local_installers/cuda-repo-ubuntu1604-9-2-local_9.2.88-1_amd64 [following]
--2022-11-06 15:24:58- https://developer.download.nvidia.com/compute/cuda/9.2/secure/Prod/local_installers/cuda-repo-ubuntu1604-9-2-local_9.2.88-1_amd64
Reusing existing connection to developer.nvidia.com:443.
HTTP request sent, awaiting response... 200 OK
Length: 1267391958 (1.2G) [application/x-deb]
Saving to: 'cuda-repo-ubuntu1604-9-2-local_9.2.88-1_amd64.deb'

cuda-repo-ubuntu1604 100%[=====] 1.18G 241MB/s in 4.7s

2022-11-06 15:25:03 (260 MB/s) - "cuda-repo-ubuntu1604-9-2-local_9.2.88-1_amd64.deb" saved [1267391958/1267391958]

Selecting previously unselected package cuda-repo-ubuntu1604-9-2-local.
(Reading database ... 115714 files and directories currently installed.)
Preparing to unpack cuda-repo-ubuntu1604-9-2-local_9.2.88-1_amd64.deb ...
Unpacking cuda-repo-ubuntu1604-9-2-local (9.2.88-1) ...
Setting up cuda-repo-ubuntu1604-9-2-local (9.2.88-1) ...

```

2m 28s completed at 8:57 PM

27°C Haze 2057 ENG IN 06-11-2022

To check CUDA installation:

**!nvcc --version**

Run the given command to install a small extension to run nvcc from the Notebook cells.

**!pip install git+https://github.com/andreinechaev/nvcc4jupyter.git**

```

[ ] libcuda1-396 : Depends: nvidia-396 (>= 396.26) but it is not going to be installed
nvidia-396-dev : Depends: nvidia-396 (>= 396.26) but it is not going to be installed
nvidia-opencl-icd-396 : Depends: nvidia-396 (>= 396.26) but it is not going to be installed
E: Unmet dependencies. Try 'apt --fix-broken install' with no packages (or specify a solution).

[1] In nvcc --version

nvcc: NVIDIA (R) cuda compiler driver
Copyright (c) 2005-2021 NVIDIA Corporation
Built on Sun_Feb_14_21:12:58_PST_2021
Cuda compilation tools, release 11.2, V11.2.152
Build cuda_11.2.r11.2/compiler.29618528_0

[1] pip install git+https://github.com/andreinechaev/nvcc4jupyter.git

Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Collecting git+https://github.com/andreinechaev/nvcc4jupyter.git
  Cloning https://github.com/andreinechaev/nvcc4jupyter.git to /tmp/pip-req-build-efz8ucn7
    Running command git clone -q https://github.com/andreinechaev/nvcc4jupyter.git /tmp/pip-req-build-efz8ucn7
Building wheels for collected packages: NVCCPlugin
  Building wheel for NVCCPlugin (setup.py): ... done
    Creating /tmp/pip-req-build-efz8ucn7/NVCCPlugin: filename=NVCCPlugin-0.0.2-py3-none-any.whl size=4306 sha256=519cae0ec89cbecacf2941fbe545d2a28a722beaf42888a6e5784fbdb0ac619
    Stored in directory: /tmp/pip-req-build-efz8ucn7/wheels/ca/33/8d/3c86eb85e97d2b6169d95c6e8f2c297fdec60db6e84cb56f5e
Successfully built NVCCPlugin
Installing collected packages: NVCCPlugin
Successfully installed NVCCPlugin-0.0.2

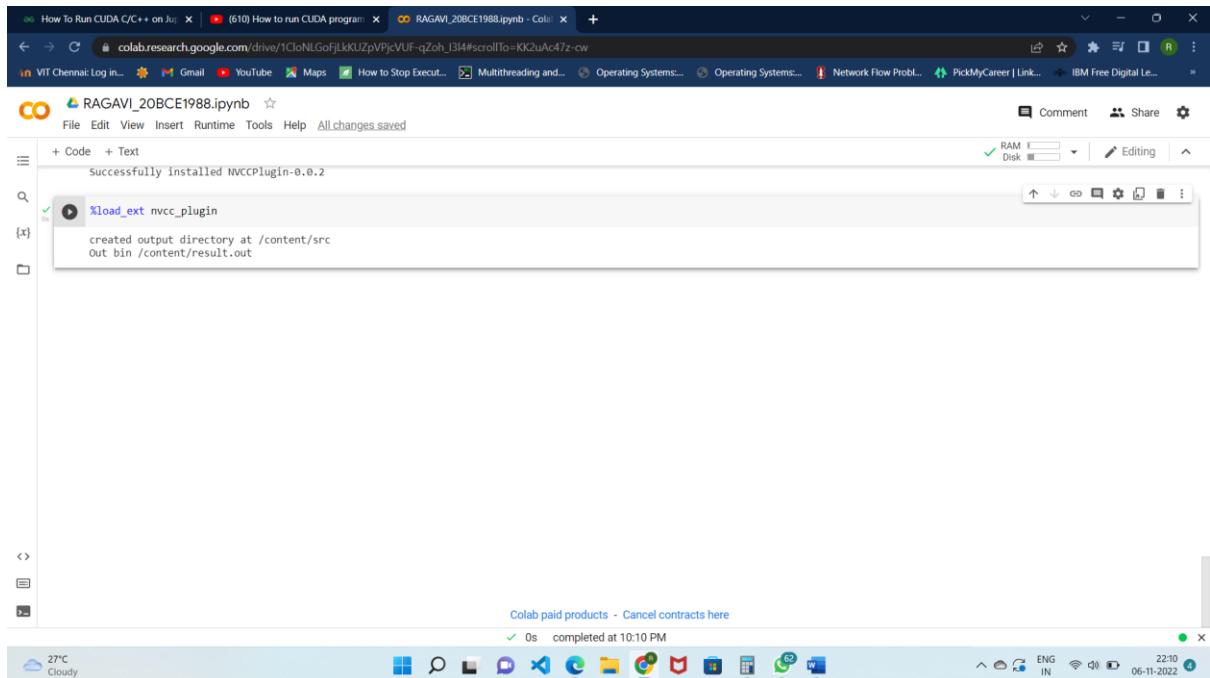
```

Automatic saving failed. This file was updated remotely or in another tab. Show diff 5s completed at 10:09 PM

27°C Cloudy 22:09 ENG IN 06-11-2022

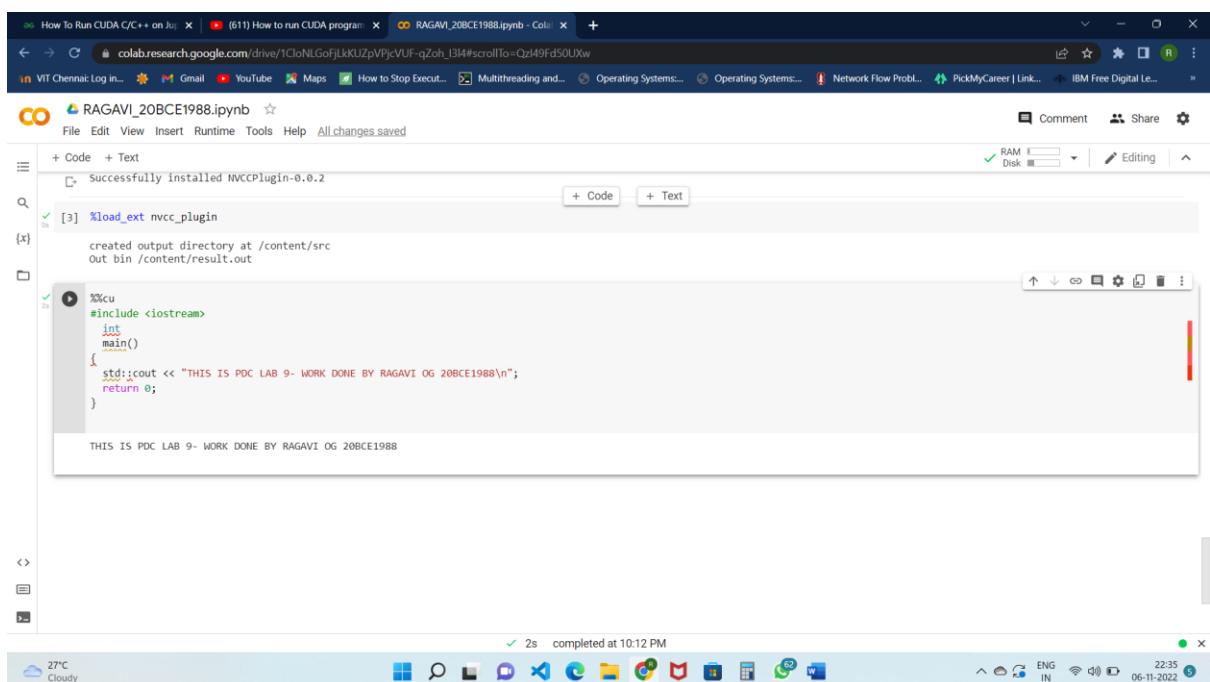
Load the extension using the code given below:

```
%load_ext nvcc_plugin
```



The screenshot shows a Jupyter Notebook interface in Google Colab. The code cell contains the command `%load_ext nvcc_plugin`. The output cell shows the message "Successfully installed NVCCPlugin-0.0.2" followed by "created output directory at /content/src". The notebook has tabs for "Code" and "Text". The status bar at the bottom indicates "0s completed at 10:10 PM".

Sample output:



The screenshot shows a Jupyter Notebook interface in Google Colab. The code cell contains a C++ program with the code 

```
%cu
#include <iostream>
int
main()
{
    std::cout << "THIS IS PDC LAB 9- WORK DONE BY RAGAVI OG 20BCE1988\n";
    return 0;
}
```

 The output cell displays the text "THIS IS PDC LAB 9- WORK DONE BY RAGAVI OG 20BCE1988". The status bar at the bottom indicates "2s completed at 10:12 PM".

**Problem: Addition of 2 integer arrays:**

## CODE:

```
%%cu
#include<stdio.h>
#include<cuda.h>

__global__ void addr(int *x,int *y, int *z)
{
    int id=blockIdx.x;
    z[id]=x[id]+y[id];
}

int main()
{
    int a[5];
    int b[5];
    int c[5];
    int *d,*e,*f;
    int i;

    for(i=0;i<5;i++)
    {
        a[i] = rand() % 100;

    }

    for(i=0;i<5;i++)
    {
        b[i] = rand() % 100;

    }

    cudaMalloc((void **) &d,5*sizeof(int));
    cudaMalloc((void **) &e,5*sizeof(int));
    cudaMalloc((void **) &f,5*sizeof(int));

    cudaMemcpy(d,a,5*sizeof(int),cudaMemcpyHostToDevice);
    cudaMemcpy(e,b,5*sizeof(int),cudaMemcpyHostToDevice);

    addr<<<5,1>>>(d,e,f);

    cudaMemcpy(c,f,5*sizeof(int),cudaMemcpyDeviceToHost);
```

```

printf("\nSum of two arrays:\n ");
for(i=0;i<5;i++)
{
    printf("%d\t",c[i]);
}

cudaFree(d);
cudaFree(e);
cudaFree(f);

return 0;
}

```

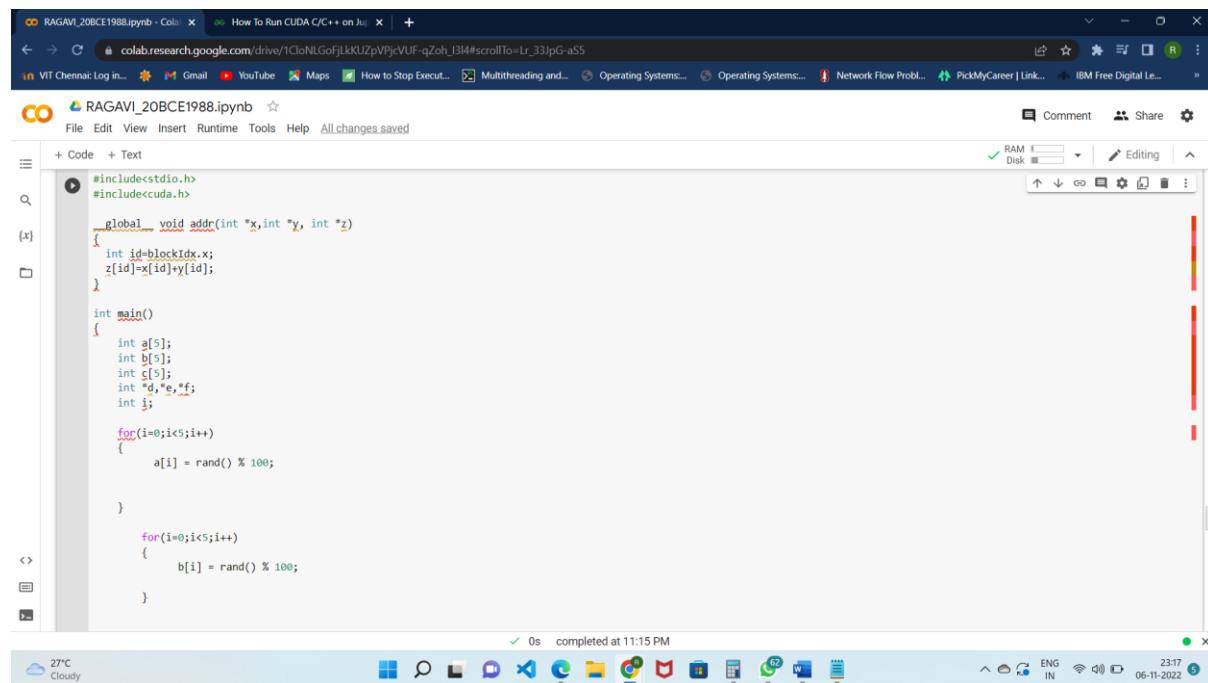
**OUTPUT:**

```

Sum of two arrays:
5 18 2 56 54

```

**ACTUAL CODE SNAPSHTOS:**



The screenshot shows a Google Colab notebook titled "RAGAVI\_20BCE1988.ipynb". The code in the cell is as follows:

```

#include<stdio.h>
#include<cuda.h>

__global__ void add(int *x, int *y, int *z)
{
    int id=blockIdx.x;
    z[id]=x[id]+y[id];
}

int main()
{
    int a[5];
    int b[5];
    int c[5];
    int d,*e,*f;
    int i;

    for(i=0;i<5;i++)
    {
        a[i] = rand() % 100;
    }

    for(i=0;i<5;i++)
    {
        b[i] = rand() % 100;
    }
}

```

The code is intended to sum two arrays (a and b) and store the result in array c. The arrays are 5x1. The code uses CUDA's global memory and block indexing.

The screenshot shows a Google Colab notebook titled "RAGAVI\_20BCE1988.ipynb" running on a Windows desktop. The code in the cell is a CUDA C program that generates random numbers for array b, allocates memory for arrays d, e, and f, copies data from host to device, calculates the sum of arrays e and f, and then prints the result. The code is annotated with comments explaining its purpose.

```
b[i] = rand() % 100;
}

{x}
cudaMalloc((void **)d,5*sizeof(int));
cudaMalloc((void **)e,5*sizeof(int));
cudaMalloc((void **)f,5*sizeof(int));

cudaMemcpy(d,a,5*sizeof(int),cudaMemcpyHostToDevice);
cudaMemcpy(e,b,5*sizeof(int),cudaMemcpyHostToDevice);

addr<<<5,1>>>(d,e,f);

cudaMemcpy(c,f,5*sizeof(int),cudaMemcpyDeviceToHost);

printf("\nsum of two arrays:\n ");
for(i=0;i<5;i++)
{
    printf("%d\t",c[i]);
}

cudaFree(d);
cudaFree(e);
cudaFree(f);

return 0;
}
```